BClib$^{\text{TM}}$: Bluespec libraries for Convey$^{\text{TM}}$ HC and MX computers

For developing HW-SW Applications using Bluespec's BSV

Release 2014.12.C

December 21, 2014

Trademarks and copyrights

Verilog is a trademark of IEEE (the Institute of Electrical and Electronics Engineers). The Verilog standard is copyrighted, owned and maintained by IEEE.

VHDL is a trademark of IEEE (the Institute of Electrical and Electronics Engineers). The VHDL standard is copyrighted, owned and maintained by IEEE.

SystemVerilog is a trademark of IEEE. The SystemVerilog standard is owned and maintained by IEEE.

Bluespec, Bluesim, AzureIp and BClib are trademarks of Bluespec, Inc.

"Convey", and other terms herein describing Convey artefacts, are trademarks of Convey Computer Corp.

Keywords and phrases: FPGA acceleration, HPC, High Performance Computing, Convey, HLS, High-Level Synthesis, PDK, Personality Development Kit, Bluespec, BSV, Bluespec SystemVerilog.

Questions or support: `support@bluespec.com`

Contents

Common abbreviations, acronyms and terminology used in this document

HW	Hardware
SW	Software
BSV	Bluespec's Bluespec SystemVerilog language
app	Application Program
HC	Generic: Convey's "Hybrid Computers"
	Specific: early generations of Convey computers had model names HC-1, HC-1ex, HC-2, HC-2ex
MX	Model name for a later generation of Convey commputers
FPGA id	The unique identity (0..3) of each of the 4 FPGAs on an HC
PDK	Convey's "Personality Development Kit" for HCs
MC, memory bank, memory port	Used interchangeably to refer to one of 8 memory controllers on an HC or one of 8 memory interfaces on an FPGA
pthread	Standard POSIX "threads" mechanism in C/C++

1   Introduction

Convey Computer Corp. (www.conveycomputer.com) offers an innovative platform for High-Performance Computing (HPC). A Convey HC "hybrid computer" consists of an industry-standard Intel x86 server, along with an FPGA-based coprocessor. The features of the coprocessor include:

- Multiple user-programmable FPGAs (four Xilinx Virtex 5 or Virtex 6 FPGAs in current HC and MX models)
- A parallel, pipelined, high-bandwidth memory system with ports on each FPGA
- "Shared memory" support between x86 server memory and FPGA memory, i.e., they share a common virtual address space
- Direct FPGA-to-FPGA communication channels
- Verilog APIs on the FPGAs to access resources (memory, communication, etc.)

Technically, the FPGA programs execute x86 co-processor instructions (although this fact is hidden in BClib). Hence, Convey calls each such FPGA programming a "personality". Convey supplies some pre-built personalities, but users can also construct their own, tuned for their applications. The user's FPGA programming does not talk directly to FPGA pins; rather, they talk to a Verilog API provided by Convey as part of their "Personality Development Kit" (PDK). User-programming of an FPGA personality may be done in the traditional way in RTL (Verilog/VHDL), or with High-Level Synthesis (HLS) tools for FPGAs.

1.1   This document, and its intended audience

This document is intended for people who wish to produce complete applications for the Convey platform, using C (or C++) for the SW side and the BSV high-level hardware design language for the HW side. In addition to using high-level BSV instead of Verilog or VHDL to design the HW side, BClib greatly simplifies SW-HW invocation and communication. BClib also enables fast, whole-application simulation (C + BSV) for testing and debugging before deployment on the hardware. Thus HW-SW development can be done completely within the facilities of BClib.

To build apps for Convey HCs using BClib, users must know C/C++ (for coding the SW) and BSV (for coding the HW), and how to use the Bluespec compiler *bsc* to build a Bluesim simulation and generate Verilog (this document shows examples of these uses). It is not necessary to be familiar with Convey's PDK (Personality Development Kit), even though BClib is built on top of Convey's PDK (in particular, the final stages of building and packaging the FPGA bitfile for deployment on the Convey platform is done using PDK facilities). We provide standalone instructions for this step in this document.

This document is not a tutorial on BSV or Bluespec tools, for which other resources are available ([1, 2] and more at www.bluespec.com). However, the code example in Sec. 3 can be used as a BSV learning exercise.

Below, you will see that we have drastically simplified the HW-SW linkage mechanism, compared to the more detailed and extensive facilities in Convey's PDK (ref. "dispatch interface" and "scalar coprocessor"). We expect that this will be adequate for the vast majority of apps written using BClib. For more advanced users who wish to use their own linkage mechanisms, BClib actually gives you full access to all the Convey PDK interfaces (see reference manual in Appendix D).

1.2 Disclaimers, licensing and (non-)support

BClib is offered by Bluespec, Inc. as an independent "after-market" library for certain hybrid computer platforms (host + FPGAs) from Convey Computer Corp. Specifically, BClib and Bluespec, Inc. have no official relationship with Convey Computer Corp. Any description in this document of Convey tools, software, development steps etc. should not be construed as the opinion of, or commitment by, Convey Computer Corp. The descriptions here merely reflect Bluespec's unofficial understanding of Convey tools and development flow, based on reading Convey's published manuals and documents.

See the `LICENSE` file in the distribution for license text. Briefly: BClib is free and open-source for anyone to use, modify and develop. Bluespec does not make any commitment to support, bug-fix, or further evolve BClib (this may happen, but solely at Bluespec, Inc.'s discretion.) BClib is provided as-is, without warranty of any kind, express or implied.

1.3 Contents of this document

Sec. 2 describes an overview of HW-SW programs developed using BClib, how they are simulated in Bluesim, and how they are deployed and executed on Convey HCs.

Sec. 3 describes in detail three versions of a small example whose complete source codes are included in the BClib distribution. Specifically, it is a BSV version of the same "vadd" (vector add) example in the Convey PDK distribution, and which is described in the Convey PDK Reference Manual. This section has a detailed code walk-through of the BSV and C codes for Version 0, and shows transcripts from building and executing it, both in Bluesim (standalone test) and on Convey HW. It also briefly describes Version 1 and Version 2, focusing on the differences from Version 0 (they only differ in the BSV source codes; the C source code and the build and execute procedures are the same).

Appendix A briefly reprises information from Convey PDK documentation on how to deploy your application on Convey HC hardware (compiling and linking the software side; converting the hardware side into a bitfile and installing it as a personality; running your application).

Appendix B briefly reprises information from Convey PDK documentation on how to run your application in the Convey PDK simulation environment (using Verilog simulation). This information is provided for completeness; we typically only simulate in Bluesim.

Appendix C has a few performance tips.

Appendix D contains a reference manual for all the facilities available in BClib.

Appendix E has brief notes on each release of this library.

2 Overview of HW-SW structure and development flow
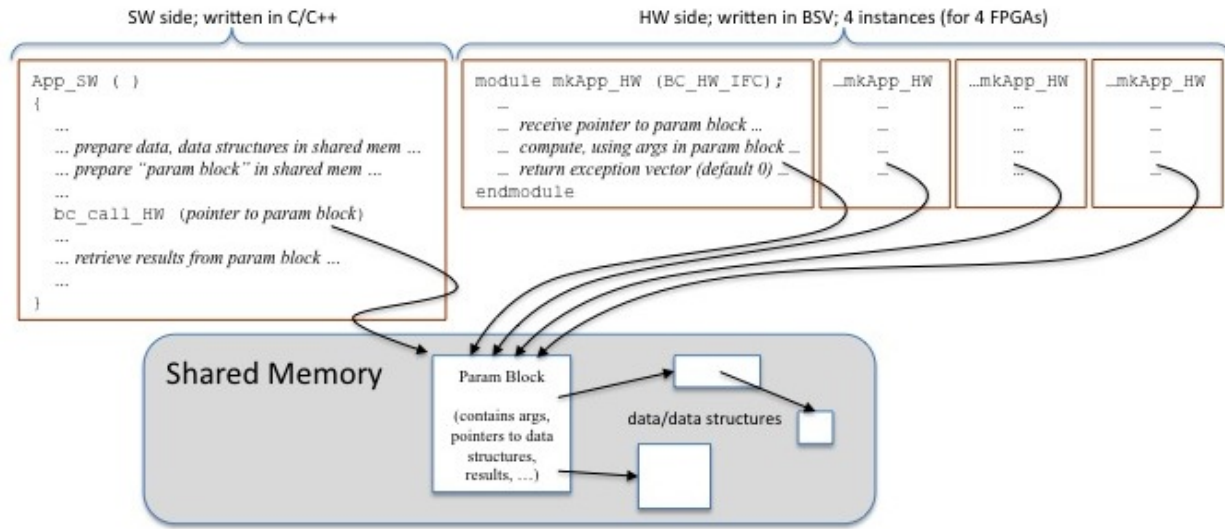
2.1 Overview: HW-SW structure

Figure 1: Overall structure of SW and HW

Fig. 1 shows the overall HW-SW structure of an app. On the left we see the SW side with the top-level function `App_SW()` which can be written in C, C++, or any such linkable language. The C/C++ code allocates and initializes any data and data structures needed for the HW side, in shared memory. It then assembles a "parameter block" in shared memory containing any information the HW side might need for its computation such as input data, pointers to data structures, pointers to data structures that will hold output data, etc. It then performs the call:

```
    bc_call_HW (pointer_to_param_block);
```

On the HW side, the BSV code receives this pointer, using which it can access all the data structures in shared memory and perform its computation, and write outputs into the param block or data structures. At the end of the HW computation, the SW thread resumes execution, i.e., the `bc_call_HW` behaves like a true procedure call to a HW routine.

Your top-level function `App_SW()` will either be called by the conventional C top-level function `main()` when running on the HC (see Sec. 2.3 below), or it will be called by `bc_SW_main()` when running in Bluesim (see Sec. 2.2 below). Thus, the latter two functions are typically just thin wrappers to launch `App_SW()`, in which the actual app work is done.

The "param block" in memory is a general argument-result passing mechanism—each app will choose the exact structure and interpretation of the param block it needs.

The HW side typically involves all 4 FPGAs in the HC. All of them receive the same param block address, but because they each know their FPGA identity (0,1,2,3), it is easy for them to have distinct arguments and results by accessing different parts of the param block based on their FPGA id.

In some applications, the SW and HW execute concurrently, i.e., it is not just a "call-return" structure, and the SW and HW are genuine coroutines. This is accommodated by having multiple pthreads on the SW side, which can be created within `App_SW()`. One thread may sit blocked in the `bc_call_HW` call, while other threads stream data to/from the HW or do other useful work. Data is communicated between SW and HW via shared memory. [We expect a future release of BClib to include libraries that will assist in building such streaming communications.]

2.2   Overview: Simulation of SW with HW using Bluesim

The structure of Fig. 1 can be completely simulated in Bluespec's Bluesim environment. In your C code, you define a BClib-standard pthread entry-point that will call your `App_SW()` function:

```
void *bc_SW_main (void * ignored);
```

This could be as simple as a one-line function calling `App_SW()`, or it may collect arguments from parameter files or environment variables before doing the call, etc.

BClib includes a top-level BSV module `mkBC_Sim_Main`. It instantiates four copies of your `mkApp_HW` top-level BSV module, and connects them to transactors that access C memory. It also creates a pthread starting at your `bc_SW_main()` function (which, in turn, will call your `App_SW()` function). The overall structure is illustrated in Fig. 2.
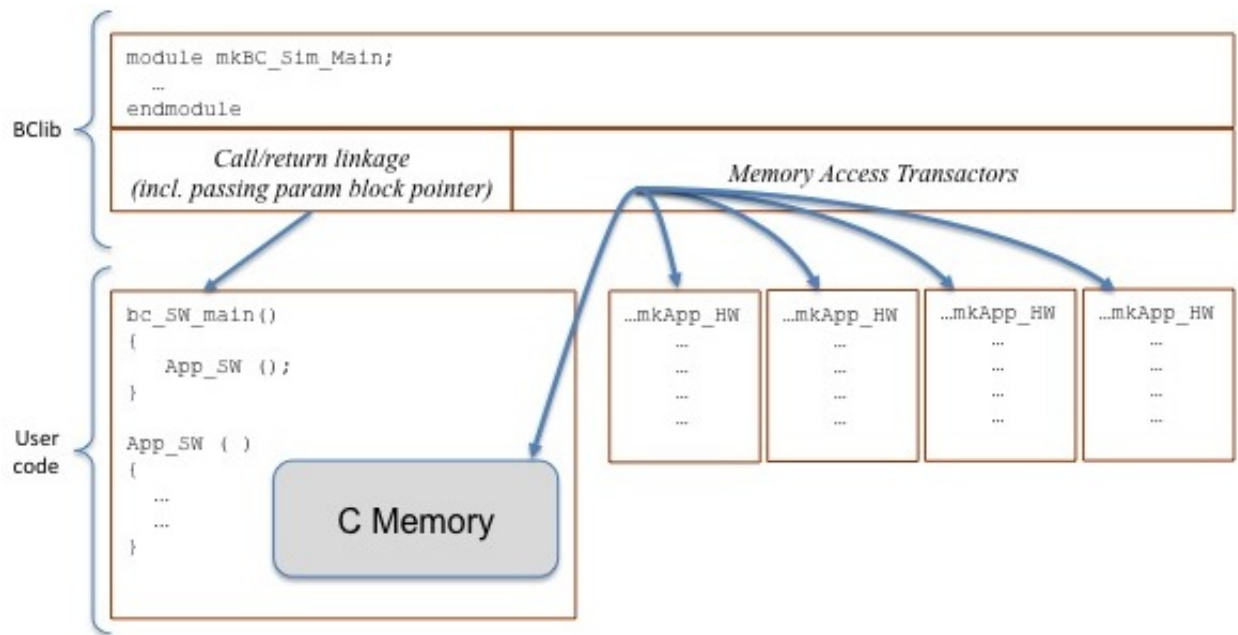
Figure 2: Structure in Bluesim

To summarize, all you have to do is:

- Create your `bc_SW_main()` entry point, which calls your `App_SW()` function.
- Do a standard *bsc* compile/link/simulate with `mkBC_Sim_Main` as the top-level module. Memory accesses in your BSV code will be serviced using C memory, i.e., "shared memory"

(More details are shown in the examples later).

2.3   Overview: Execution of SW with HW on Convey HC

After checking that your app is working Bluesim simulation, you're ready to deploy it on the HC. For the HW side, the steps are:

- Use *bsc* to compile the BClib-supplied BSV file `Cae_Pers.bsv` and its dependencies into Verilog. Since it instantiates your `mkApp_HW` module, that will be compiled as well, along with its dependencies. The top-level generated Verilog module is `cae_pers.v`, as required by the Convey PDK.

- Within the Convey PDK environment, use the PDK's `make` command to synthesize this Verilog into an FPGA bitfile. This will also link in the necessary Convey PDK Verilog files. It will invoke the Xilinx synthesis tools in your installation. As is usual for FPGA synthesis, you may have to iterate to this step several times if you do not meet timing (the default timing target is 150 MHz).
- Use the PDK's `make release` command to create a standard PDK release tar file.
- Install the tar file as an available "personality" on the HC (details later, but this just involves copying the tar file to a standard directory on the HC, and possibly fixing up a symbolic link to it).

For the SW side, the steps are:

- Your C source files should now include the traditional `main()` function, and it should invoke your `App_SW()` function. This could be as simple as a one-line function calling `App_SW()`, or it may collect arguments from the command line, environment variables, or parameter files before doing the call, etc. Your C source files may also include calls to Convey PDK runtime libs that provide versions of malloc(), memcpy() etc. so that you can "place" data structures close to the FPGAs or close to the x86 on the HC (we typically "ifdef" these so that when compiling for Bluesim they revert to ordinary malloc(), memcpy(), etc.)
- Compile your C sources using Convey's C compiler (similar to gcc).
- Compile the BClib-supplied `cpAsm.s` file using Convey's C compiler. This is a boilerplate assembly language program for the scalar coprocessor that participates in linking your C code to your BSV code.
- Link the object files just created using Convey's C compiler. This will also link in Convey PDK run-time libraries.

Now, execute your program on the HC, just like you would execute any C-compiled executable. The overall structure is shown if Fig. 3.
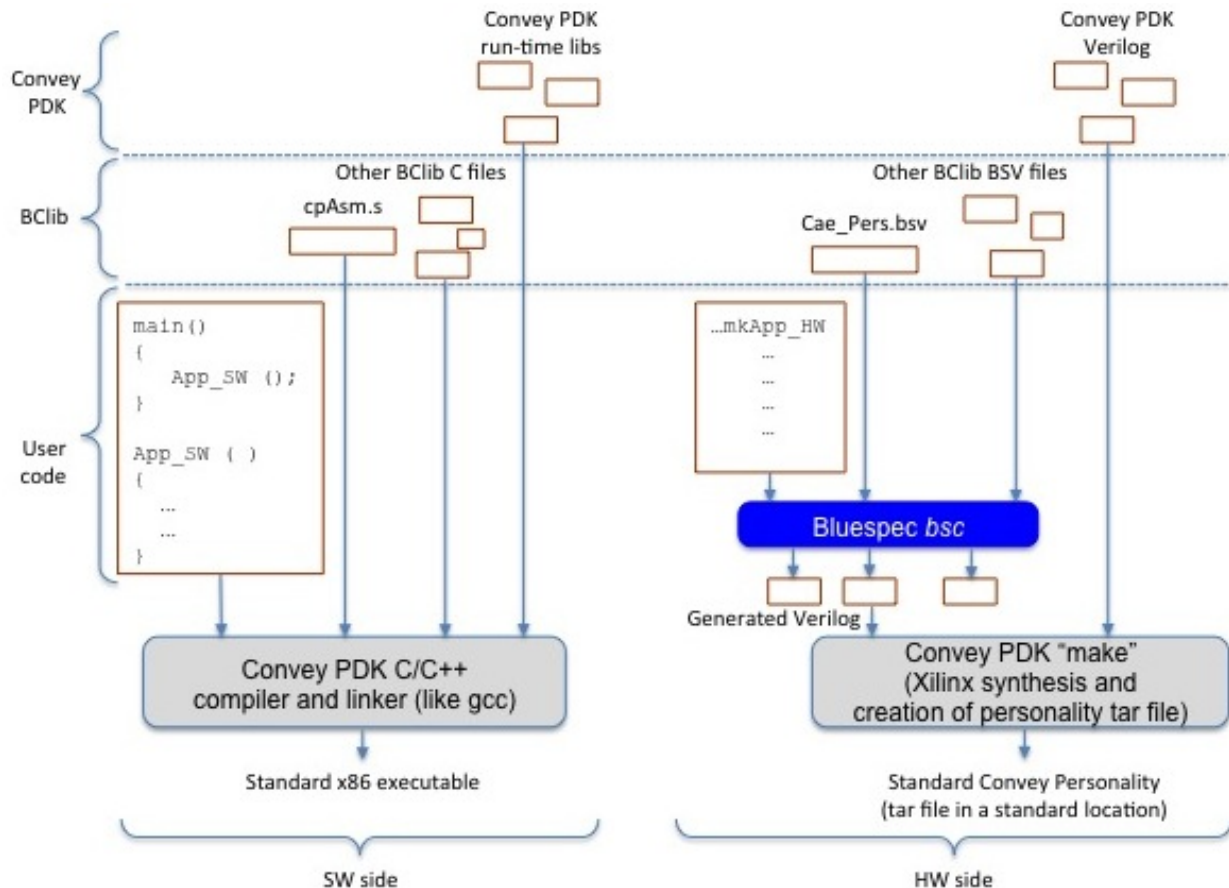

Figure 3: Structure for HC execution

(More details are shown in the examples later).

3 Fully-worked tutorial examples

3.1 Problem statement

In this section we describe three tutorial examples in detail—source code walk-throughs, simulation, and execution on the HC. The source codes are included in the BClib distribution in the `Example` sub-directory. These are three versions of the same "vector add" example that is found in Convey's PDK. Functionally it is very simple: it performs an element-by-element addition of two vectors V1 and V2 into an output vector V3, and also sums up V3:

$$\forall_{j=0}^{n-1} \quad V3_j = V1_j + V2_j$$

$$sum = \sum_{j=1}^{n-1} V3_j$$

The input vectors are initialized as follows:

$$\forall_{j=0}^{n-1} \quad V1_j = j \quad \text{and} \quad V2_j = 2j$$

For example, for vectors of 100 elements ($n = 100$), the expected sum is 14850. In these codes, all the data elements are 64-bit integers.

The essential differences between the three versions are the following:

- **Version 0** is nearly identical to the Convey PDK Verilog example. It assumes a simple static memory mapping—all three vectors are assumed to be aligned identically with respect to HC's 8 memory controllers (MCs), so that it is easy to direct the two reads and writes for each addition to the correct MC. Further, it assumes memory read responses arrive in the same order as the corresponding requests.

- **Version 1** also assumes the same simple static memory mapping, but does not assume in-order memory read responses. The BSV code manages memory re-ordering.

- **Version 2** does not assume any particular alignment of the vectors with respect to the MCs. Thus, it contains a "router" that directs reads and writes to the correct MC, and it also internally manages memory re-ordering.

Why is this interesting? Briefly (more details in the reference section), the Convey HC's memory is organized into 8 banks, each with its own memory controller. Bytes are "striped" across banks: an address a is located in bank $a[8:6]$. This is exposed to the FPGAs, i.e., each FPGA has 8 memory ports connecting to the banks, which can be operated in parallel. It is the responsibility of the programming on each FPGA to direct each memory request to the correct port (bits [8:6]). A request sent to the "wrong" port will simply return the value that is aliased to the same address modulo bits [8:6], i.e., likely a wrong value.

The "leanest" applications (highest peformance, least FPGA resources, simplest structure) will likely have static data structure mappings so that memory requests are sent directly to the correct MC. But in the majority of applications such static mappings are not feasible; these will need mechanisms in the FPGA programming to route requests and responses to the correct MC.

3.2   Vector-add example, version 0

We assume that Convey HC memory is configured with the simple "binary interleave", as described in Sec. D.2.1 and illustrated in Fig. 5. We refer to each 512B (byte) row as a "stripe" across the 8 MCs (Memory Controllers). Bits [8:0] of an address represent its offset within a stripe, bits [8:6] identify the MC it belongs to, and bits [5:0] represent its offset within the 64B chunk in that MC.

Fig. 4 shows how the three vectors are laid out in this binary interleave. Specifically, note that we assume

Figure 4: Memory layout of the three vectors in version 1 of the vector-add example

that all three vectors have the same starting offset within their respective stripes (version 2 will relax this constraint). The figure shows V1, V2 and V3 at increasing addresses, but this is not assumed. We only assume that V3 does not overlap V1 or V2 such that writes into V3 may disturb V1 and V2.

The SW side (C code) sets up the computation by initializing the vectors and computing the expected sum. It then creates a parameter block containing, for each FPGA,

- the size of the segments of V1, V2 and V3 for which that FPGA is responsible (roughly $\frac{1}{4}$ of the vectors),
- the starting addresses of its segments of V1, V2 and V3, and
- slots where the FPGA will write its final partial sum.

It then performs the `bc_call_HW()` call, passing it the argument of the param block.

The HW side (BSV code) receives this pointer to the param block, and performs its computation using several levels of parallelism:

- Each of the 4 FPGAs independently computes the partial sum of roughly $\frac{1}{4}$ of the vectors (we round up each segment to a 512-Byte boundary (= 64 64-bit words), so that it is aligned with MC 0). These partial sums are combined on the SW side (by the C code) to produce the full sum.

- Within each FPGA, the code contains 16 parallel "channel adders" (8 pairs), each pair handling parts of the vectors that lie within one MC. Each engine does sequential traversals of its pieces of the vectors, as illustrated by the red squiggly arrows in MC5 in Fig. 4. "Even" engines handle even elements of the vectors and use the even port of each even/odd pair of MC ports, and "odd" engines handle odd elements using odd ports. At each point, the engine performs the + operation, and accumulates a partial sum.

  At the end, the partial sums from the channel adders are combined to give the partial sum for the FPGA.

- Within each channel adder, memory traffic is parallel and pipelined. A generator process streams out memory read requests for data from V1 and V2. A second, independent process receives these pairs of results as they arrive (which can be 100s of cycles later), performs the addition and partial sum, and streams out write requests to V3.

In steady state, on each clock, each channel adder should issue a memory request to its port. Since each vector element involves 3 memory requests (2 reads, 1 write), in the best case we should take 3N cycles to compute the result for a vector of size N (the best case is when every cycle of the memory port is occupied with a result).

Performance note: On the HC-1 platform we measured a performance penalty if read and write requests were interleaved at fine granularity, making this app initially 40% slower than the Convey PDK example vector-add (which is written in Verilog). `App_HW_v0.bsv`, therefore, alternates bursts of 1536×2 reads with bursts of 1536 writes, and performances matches that of the the PDK example.

On each FPGA, the HW computation writes its final partial sum into its slot in the param block, before completing. The C code then returns from its `bc_call_HW()` call. It then reads the four FPGA-specific partial sums from the param block, and sums them for the final sum, and checks it against the expected final sum.

For simplicity in these examples, all our data (including vector elements and parameter block elements) are 8-Byte (64-bit values). The Convey HC is capable of accessing memory at 1, 2, 4 and 8-Byte granularities.

3.2.1   SW code walk-through

Please refer to the code which is in the file: `Example/C_src/App_SW.c`
This is about 220 lines of C.

At the bottom of the file you will see the traditional `main()` routine, which is the entry point when running on the actual Convey HC. It is enclosed in "`#ifdef FOR_HW`" because it is not used when executing on Bluesim. It is just a thin wrapper containing lines like this::

```
                                      App_SW.c
int main (int argc, char *argv[])
{
    ...
    bc_init_HW ("65123");
    ...
        return App_SW (argv [1]);
}
```

It first calls a BClib initialization function. The character string argument is the "signature name" chosen for this Convey HW personality (see A for more details). Then, it calls `App_SW()`, where the real work is done.

Just preceding `main()` is this function:

```
                                      App_SW.c
void *bc_SW_main (void * ignored)
{
    ...
    App_SW (s);
    ...
}
```

which is the pthread entry point when running under Bluesim. It is also a thin wrapper that just calls `App_SW()`. You can see in the code for `main()` that it picks up the desired vector size (number of elements) from the command line argument, if given. Both `main()` and `bc_SW_main()` also look for the vector size in the environment variable `VECTOR_SIZE`. If the size is not given, it defaults to 100.

Above these is the common function used both in Bluesim and in HC execution:

```
──────────────────────────── App_SW.c ────────────────────────────
int App_SW (const char *vsize_s)
{
    ...
}
```

After some initial code to determine the desired vector size, it allocates the three vectors using statements like this:

```
──────────────────────────── App_SW.c ────────────────────────────
    vin1 = (uint64_t *) bc_malloc ("vin1", vsize * 8, align512,
                                   fpga_side, NULL, exit_on_fail);
```

The arguments provide the desired byte size, alignment (to a 512-Byte stripe), a request to allocate it on the "FPGA side" of the HC, and to fail on exit. Ordinary malloc() is used when compiling for Bluesim.

After initializing the vectors and calculating the expected sum, it prepares the parameter block which carries information to the HW side:

```
──────────────────────────── App_SW.c ────────────────────────────
    for (fpga = 0; fpga < NUM_FPGAs; fpga++) {
      uint64_t lo = fpga * vsize_per_fpga;
      uint64_t hi = min (vsize, (fpga + 1) * vsize_per_fpga);
      param_block [PARAM_VSIZE*8+fpga] = ((lo > hi) ? 0 : (hi - lo));
      param_block [PARAM_VIN1 *8+fpga] = ptr_to_ui64 (vin1 + (fpga * vsize_per_fpga));
      param_block [PARAM_VIN2 *8+fpga] = ptr_to_ui64 (vin2 + (fpga * vsize_per_fpga));
      param_block [PARAM_VOUT *8+fpga] = ptr_to_ui64 (vout + (fpga * vsize_per_fpga));
      param_block [PARAM_PARTIAL_SUM * 8 + fpga] = 0;
    }
```

Note that it fills in the vector size and starting addresses of the V1, V2 and V3 segments for each FPGA. The convenience function `ptr_to_ui64` is provided in BClib to convert pointers to 64-bit values. The param block also has slots for the per-FPGA partial sums. These are initialized to 0 in the last line above, but that is not strictly necessary as each FPGA will overwrite these with its partial sum.

The code then "calls" the HW routine, passing it a pointer to the param block:

```
──────────────────────────── App_SW.c ────────────────────────────
    bc_call_HW (ptr_to_ui64 (param_block));
```

In the actual source file this is surrounded by some library calls to measure and print elapsed time. When it returns from the call, the per-FPGA partial sums will have been written into the param block. The C code collects these to compute the final sum, and checks it against the expected sum:

```
──────────────────────────── App_SW.c ────────────────────────────
    actual_sum = 0;
    for (fpga = 0; fpga < NUM_FPGAs; fpga++) {
```

```
        uint64_t partial_sum  = param_block [PARAM_PARTIAL_SUM * 8 + fpga];
        printf ("C: partial sum [%0d] = %0lld\n", fpga, partial_sum);
        actual_sum += partial_sum;
    }
    printf ("C: actual sum = %0lld\n", actual_sum);

    if (expected_sum == actual_sum)
        printf ("C: PASSED (actual = expected)\n");
    else {
        printf ("C: FAILED - expected sum = %lld\n", expected_sum);
        return 1;
    }
```

3.2.2   HW code walk-through

Please refer to the code which is in the file: Example/BSV_src/App_HW_v0.bsv
This is about 565 lines of copiusly-commented BSV. This walk-through is fairly detailed, to provide a brief
tutorial for those new to BSV. The overall structure of the file looks like this:

App_HW_v0.bsv

```
package App_HW;

// ================================================================
// This is Version 0 of the Vector-Add example ...
...
... various library imports
...
// ================================================================
// Top-level of the app (for 1 FPGA)

module mkApp_HW (BC_HW_IFC);
    ...
endmodule

// ================================================================
// HW module with simplified API
...
module mkApp_HW2 (BC_HW2_IFC);
    ...
endmodule

// ================================================================
// The following module is instantiated 16 times
//     once for each MC port
...
module mkChanVadd (ChanVadd_IFC);
    ...
endmodule

endpackage
```

It defines a BSV "package" (a namespace), imports some standard BSV libraries, imports the BClib packages `BC_HW_IFC` and `BC_Transactors` and defines three modules. The first module, `mkApp_HW`, is the BClib-prescribed top-level module with the prescribed `BC_HW_IFC` interface. It instantiates the second module, `mkApp_HW2`. This, in turn, creates 16 instances of the third module `mkChanVadd`.

The top-level module `mkApp_HW` is mostly boilerplate. In fact it is identical to the corresponding `mkApp_HW` in v2 of this example (file `Examples/BSV_src/App_HW_v2.bsv`). If you compare this module to the corresponding `mkApp_HW` in v1 of this example (file `Examples/BSV_src/App_HW_v1.bsv`), you will see that the only difference is that v1 uses an "ordering" memory transactor, whereas v0 and v2 do not. It is likely that most applications you write will be able to use one or the other of these two boilerplate modules as is. They will change only if you need more full-bore access to Convey interfaces (Dispatch, Management). It contains an FSM that is an infinite loop that waits for a "caep" instruction (the "start" signal from the SW), starts the computation by calling `hw2.start()`, and waits for the computation to complete, by calling `hw2.waitTillDone()` before looping back to wait for the next caep instruction.

The interface to second-level module `mkApp_HW2` is defined as follows:

```
————————————— App_HW_v0.bsv —————————————
interface BC_HW2_IFC;
    method Action start (BC_AEId fpga_id, BC_Data param_block_addr);
    method Action waitTillDone;

    interface Vector #(16, BC_MC_Client) mc_ifcs;
endinterface: BC_HW2_IFC
```

The parent module calls the `start` method to begin the HW computation, and calls the `waitTillDone` method to await completion. During the computation, traffic to and from memory flows through the `mc_ifcs` interfaces (requests and responses). This is a vector of 16 "memory clients", one per Convey PDK-supplied memory port. The definition of `BC_MC_Client` is in the BClib file `BC_HW_IFC.bsv`.

The contents of the module are application specific. The parameter block communicated from SW to HW is an array of 64-bit values, laid out as follows for simplicity (recall that each FPGA is responsible for about $\frac{1}{4}$ of the full vectors):

- locations [0], [1], [2] and [3] (in MC 0) hold the size of the vectors that FPGAs 0, 1, 2 and 3 are responsible for
- locations [8], [9], [10] and [11] (in MC 1) hold base addresses of the segments of the first input vector that FPGAs 0, 1, 2 and 3 are responsible for
- locations [16], [17], [18] and [19] (in MC 2) hold base addresses of the segments of the second input vector that FPGAs 0, 1, 2 and 3 are responsible for
- locations [24], [25], [26] and [27] (in MC 3) hold base addresses of the segments of the output vector that FPGAs 0, 1, 2 and 3 are responsible for
- locations [32], [33], [34] and [35] (in MC 4) will hold the resulting partial sums from FPGA 0, 1, 2 and 3 at the end of the computation.

Thus each FPGA j can access its arguments from $[8 \times parameter + j]$. These parameters are defined next in the file:

```
————————————— App_HW_v0.bsv —————————————
Integer param_VSIZE       = 0;
Integer param_VIN1_P      = 1;
Integer param_VIN2_P      = 2;
Integer param_VOUT_P      = 3;
Integer param_PARTIAL_SUM = 4;
```

The input parameters (size and base addresses) are calculated and written to the parameter block by the C side before invoking the hardware. Each FPGA reads its arguments, performs its computation, and writes its partial sum back as the output parameter. These are then read by the C side. Next, we have:

```
──────────── App_HW_v0.bsv ────────────
(* synthesize *)
module mkApp_HW2 (BC_HW2_IFC);
   Reg #(BC_AEId)        rg_fpga_id          <- mkRegU;
   Reg #(BC_Addr)        rg_param_block_addr <- mkRegU;
   Reg #(BC_Data)        rg_partial_sum      <- mkRegU;
```

The (* synthesize *) attribute tells *bsc*, the Bluespec compiler, to create a separate Verilog module for this BSV module (otherwise it would just be inlined where it is instantiated, in the parent module mkApp_HW). In the module, we instantiate 3 registers to hold the FPGA id (0..3), the address of the param block, and the partial sum for this FPGA. Next, we instantiate FIFOs for memory read/write requests and responses, and for flush requests and responses, for each of the 16 memory ports (there are a total of 64 FIFOs).

```
──────────── App_HW_v0.bsv ────────────
   // ------------------------------------------------------------
   // MC FIFOs

   Vector #(16, FIFOF #(BC_MC_REQ))       f_reqs       <- replicateM (mkFIFOF);
   Vector #(16, FIFOF #(BC_MC_RSP))       f_rsps       <- replicateM (mkFIFOF);
   Vector #(16, FIFOF #(BC_MC_flush_req)) f_flush_reqs <- replicateM (mkFIFOF);
   Vector #(16, FIFOF #(BC_MC_flush_rsp)) f_flush_rsps <- replicateM (mkFIFOF);
```

Next, we create 16 instances of the per-MC "adders", and connect their memory ports to the corresponding memory FIFOs:

```
──────────── App_HW_v0.bsv ────────────
   // ------------------------------------------------------------
   // Instantiate 16 vadders, and connect them to the MC fifos

   Vector #(16, ChanVadd_IFC) v_chanvadds <- replicateM (mkChanVadd);

   for (Integer mc = 0; mc < 16; mc = mc + 1)
      mkConnection (v_chanvadds [mc].mc_client,
                    fn_FIFOFs_to_MC_Server (f_reqs [mc],
                                            f_rsps [mc],
                                            f_flush_reqs [mc],
                                            f_flush_rsps [mc]));
   end
```

The next section has three convenience functions:

```
──────────── App_HW_v0.bsv ────────────
   function Action send_rd_req (BC_Addr base, Integer param_id);
      ...
   function Action send_wr_req (BC_Addr base, Integer param_id, BC_Data x);
      ...
```

```
    function ActionValue #(BC_Addr) recv_rd_rsp (Integer param_id);
     ...
    function Action recv_wr_rsp (Integer param_id);
    ...
```

These implement the $[8 \times parameter + j]$ address calculations for the parameters in the parameter block, enqueing requests to and dequeuing responses from the appropriate even members of the 8 MC fifo groups (we arbitrarily pick the even ports for these initial parameter reads). The `recv_wr_rsp` function is of `Action` type because it does not return any value; its only job is to consume the memory write-response from the Convey hardware, which assures us that the write has completed.

The next section is an FSM that describes the "process" executed by this module.

```
———————————————————— App_HW_v0.bsv ————————————————————
    Reg #(UInt #(5)) rg_mc <- mkRegU;


    let fsm <- mkFSM (
        seq
           ... to be described below ...
        endseq
        );
```

The register `rg_mc` is used as a loop index inside the FSM. The FSM first sends four parallel requests to memory to fetch the parameter values:

```
———————————————————— App_HW_v0.bsv ————————————————————
        // Send read requests for the parameters for this FPGA (in parallel)
        action
           send_rd_req (rg_param_block_addr, param_VSIZE);
           send_rd_req (rg_param_block_addr, param_VIN1_P);
           send_rd_req (rg_param_block_addr, param_VIN2_P);
           send_rd_req (rg_param_block_addr, param_VOUT_P);
        endaction
```

Next, in a single atomic action, it waits for, and receives the parameter values from memory (in parallel); it feeds them, in parallel, to the `start` method of all 16 per-MC adders. The base addresses are rounded up, if necessary, to the first address owned by the corresponding channel using the function `bc_base_addr_for_chan()` which is found in the file `BC_HW_IFC.bsv` It also initializes `rg_partial_sum` to 0:

```
———————————————————— App_HW_v0.bsv ————————————————————
        // Receive the parameters for this FPGA (read responses, in parallel)
        action
           let vsize   <- recv_rd_rsp (param_VSIZE);
           let vin1_p  <- recv_rd_rsp (param_VIN1_P);
           let vin2_p  <- recv_rd_rsp (param_VIN2_P);
           let vout_p  <- recv_rd_rsp (param_VOUT_P);
           for (Integer mc = 0; mc < 8; mc = mc + 1) action
               v_chanvadds [2 * mc].start      (rg_fpga_id, fromInteger (mc),
                                                 True,
                                                 bc_base_addr_for_chan (mc, vin1_p),
                                                 vin1_p + (vsize << 3),
```

```
                                                 bc_base_addr_for_chan (mc, vin2_p),
                                                 vin2_p + (vsize << 3),
                                                 bc_base_addr_for_chan (mc, vout_p),
                                                 vout_p + (vsize << 3));
             v_chanvadds [2 * mc + 1].start (rg_fpga_id, fromInteger (mc),
                                                 False,
                                                 8 + bc_base_addr_for_chan (mc, vin1_p),
                                                 vin1_p + (vsize << 3),
                                                 8 + bc_base_addr_for_chan (mc, vin2_p),
                                                 vin2_p + (vsize << 3),
                                                 8 + bc_base_addr_for_chan (mc, vout_p),
                                                 vout_p + (vsize << 3));
         endaction
         rg_partial_sum <= 0;
      endaction
```

Finally, the FSM collects and sums the per-MC partial sums results, and writes the sum back to the correct result slot in the param block:

─────────────── App_HW_v0.bsv ───────────────
```
      // Collect per-channel partial sums, accumulate into full (per-FPGA) sum
      for (rg_mc <= 0; rg_mc < 16; rg_mc <= rg_mc + 1) action
          let chan_sum <- v_chanvadds [rg_mc].result;
          rg_partial_sum <= rg_partial_sum + chan_sum;
      endaction

      // Write final (per-FPGA) sum back to param block, and drain the response
      send_wr_req (rg_param_block_addr, param_PARTIAL_SUM, rg_partial_sum);
      recv_wr_rsp (param_PARTIAL_SUM);
```

Due to BSV semantics, the FSM will automatically stall as necessary to wait for memory responses, v_chanvadds [..].results, etc.

After the FSM definition, all that remains in this module is the definition of the interface methods and sub-interfaces:

─────────────── App_HW_v0.bsv ───────────────
```
  method Action start (BC_AEId fpga_id, BC_Data param_block_addr) if (fsm.done);
      rg_fpga_id          <= fpga_id;
      rg_param_block_addr <= truncate (param_block_addr) + (extend (fpga_id) << 3);
      fsm.start;
  endmethod

  method Action waitTillDone() if (fsm.done);
      noAction;
  endmethod

  interface mc_ifcs = genWith (fn_mkMC_Client);
```

The start method is only enabled when the FSM is idle (fsm.done condition). It stores its arguments in registers and starts the FSM. The waitTillDone method is only enabled when the FSM has completed,

and otherwise performs no action. The MC interfaces just use a help function, defined earlier, to convert the FIFOs into the appropriate client interfaces.

We now discuss the last module, the individual channel-adder. Recall that its parent module `mkApp_HW2` instantiates this 16 times, once for each of the 8 even and odd MC ports. Its interface is:

```
                                      App_HW_v0.bsv
interface ChanVadd_IFC;
   method Action start (BC_AEId aeid, BC_MC chan,
                        Bool  evenNotOdd,
                        BC_Addr rd_base_1,  BC_Addr rd_limit_1,
                        BC_Addr rd_base_2,  BC_Addr rd_limit_2,
                        BC_Addr wr_base,    BC_Addr wr_limit);
   method ActionValue #(BC_Data) result ();    // chan partial sum


   interface BC_MC_Client  mc_client;
endinterface
```

The `start` method receives the FPGA id (0..3), the channel id (0..7), a boolean indicating whether this is responsible for even or odd elements of the vectors, and the base and limit addresses of the segments of the three vectors for which this FPGA is responsible. The `result` method will return the partial sum for those part of these vectors that lie in MC `chan`. The memory `mc_client_pair` interface connects to the MC for this channel.

The module `mkChanVadd` implementing this `ChanVadd_IFC` interface begins like this:

```
                                      App_HW_v0.bsv
(* synthesize *)
module mkChanVadd (ChanVadd_IFC);
   // These are only needed for $displays
   Reg #(BC_AEId)  rg_aeid <- mkRegU;
   Reg #(BC_MC)    rg_chan <- mkRegU;

   Reg #(BC_Data)          rg_partial_sum <- mkRegU;


   Reg #(RUN_STATE)        rg_run_state   <- mkReg (RUN_STATE_IDLE);


   Reg #(Bool)             rg_evenNotOdd  <- mkRegU;


   // Input vector 1
   Reg #(BC_Addr)          rg_rd_base_1   <- mkRegU;
   Reg #(BC_Addr)          rg_rd_limit_1  <- mkRegU;

   Reg #(BC_Addr)          rg_rd_base_2   <- mkRegU;
   Reg #(BC_Addr)          rg_rd_limit_2  <- mkRegU;
   Reg #(Bool)             rg_rd_active_2 <- mkReg (False);

   // When rg_reading is True, we do a run of 1536 reads from vin1 and vin2
   // When False, we do a run of 1536 writes to vout
   Reg #(Bool)             rg_reading     <- mkRegU;

   // Read sequencing
```

```
    Reg #(Bit #(16))        rg_rd_pc        <- mkRegU;
    Reg #(Bool)             rg_vin1_turn    <- mkRegU;
```

It instantiates registers to hold some of the parameters, the partial sum that it computes, a "run state" variable, and the base and limit addresses of the input vectors. The `rg_rd_active_2` register holds a boolean representing the condition that base is < than limit.

As mentioned eariler, for performance reasons we alternate bursts of 1536×2 reads and 1536 writes. The variable `rg_reading` indicates whether we are in a read burst or in a write burst. The variable `rg_rd_pc` counts off a read burst (0..1535) and, within a read burst `rg_vin1_turn` indicates whether we are reading from input vector V1 or from V2.

Next, we instantiate some variables for managing output writes:

```
───────────────────────── App_HW_v0.bsv ─────────────────────────
    // Output vector
    Reg #(BC_Addr)          rg_wr_base      <- mkRegU;
    Reg #(BC_Addr)          rg_wr_limit     <- mkRegU;
    Reg #(Bool)             rg_wr_active    <- mkReg (False);


    // Write sequencing
    Reg #(Bit #(16))        rg_wr_pc        <- mkRegU;
    Reg #(Bit #(16))        rg_wr_rsp_drain_count <- mkRegU;
```

In addition to the vector base and limit, the `rg_wr_active` register holds a boolean representing the condition that base is < than limit. `rg_wr_pc` counts off a write burst (0..1535). `rg_wr_rsp_drain_count` is a counter indicating how many writes requests have been issued, and therefore how many write responses must be consumed. Next, we instantiate FIFOs for memory requests and resonses:

```
───────────────────────── App_HW_v0.bsv ─────────────────────────
    // Mem req/rsp fifos
    FIFOF #(BC_MC_REQ)      fifo_reqs       <- mkFIFOF;
    FIFOF #(BC_MC_RSP)      fifo_rsps       <- mkFIFOF;
```

The following two registers are used to hold the response of a read from the first input vector while we are waiting for the corresponding read response from the second input vector:

```
───────────────────────── App_HW_v0.bsv ─────────────────────────
    Reg #(BC_Data)          rg_x1           <- mkRegU;
    Reg #(Bool)             rg_x1_valid     <- mkRegU;
```

The final piece of state in this module is a buffer to hold 1536 sums to be written back to memory in the next write burst:

```
───────────────────────── App_HW_v0.bsv ─────────────────────────
    // Buffer to hold 1536 results, since we alternate 1536x2 reads and 1536 writes
    // We use 'noReset' to avoid load on the Reset tree
    FIFOF #(BC_Data) fifo_sums <- mkSizedBRAMFIFOF (1536, reset_by noReset);
```

Even channel-adders start at a base address A and access the following sequence of addressess:

A+0x00	A+0x10	A+0x20	A+0x30
A+0x200+0x00	A+0x200+0x10	A+0x200+0x20	A+0x200+0x30
A+0x400+0x00	A+0x400+0x10	A+0x400+0x20	A+0x400+0x30

... and so on ...

Odd channel-adders access the following sequence of addressess:

A+0x08	A+0x18	A+0x28	A+0x38
A+0x200+0x08	A+0x200+0x18	A+0x200+0x28	A+0x200+0x38
A+0x400+0x08	A+0x400+0x18	A+0x400+0x28	A+0x400+0x38

... and so on ...

The following function, next in the code, performs the above address increments:

```
──────────────────────── App_HW_v0.bsv ────────────────────────
   function BC_Addr addr_incr (BC_Addr addr);
```

Conceptually, it increments bits [5:4] (i.e., adds 0x10), and if there is an overflow, it increments bits [47:9] (i.e., adds 0x200). Next, we have this rule:

```
──────────────────────── App_HW_v0.bsv ────────────────────────
   rule rl_init (rg_run_state == RUN_STATE_INIT);
```

which initializes various variables. Functionally, we could have eliminated this rule and merged these actions into the `start` method. Separating them into this rule is purely out of hardware considerations: we want the `start` method merely to write its arguments into registers, and all isolate inputs from all the initialization logic, such as the comparisons in the expressions in this rule. These techniques improve the clock speed at which the circuit can run.

The two rules `rl_rd_gen_0` and `rl_rd_gen_1` alternate, generating read requests to the two input vectors. The first rule is:

```
──────────────────────── App_HW_v0.bsv ────────────────────────
   rule rl_rd_gen_0 (rg_rd_active_2 && rg_vin1_turn);
      let req  = BC_MC_REQ {cmd_sub: REQ_RD, rtnctl: 0, len: BC_8B,
                            vadr: rg_rd_base_1, data: ?};

      let next_base  = addr_incr (rg_rd_base_1);
      rg_rd_base_1  <= next_base;
      fifo_reqs.enq (req);

      rg_vin1_turn    <= False;
      rg_wr_pc <= 0;
   endrule
```

The second rule is almost identical, but it also computes `rg_rd_active_2`, the termination condition that checks if we have consumed all our input elements. Note that the rules are only active while we are within the base and limit for the input vectors. The read request is issued to the current `rg_rd_base` address. The Convey `rtnctl` field is left unspecified because we are assuming memory responses to be returned in order, and so there is no need for a unique ordering tag here. In preparation for the next request, we increment the address to the next address using the `addr_incr` function: Rule `rl_gen_1` also uses `rd_pc`

to count off 1536 (pairs of) reads and, when done, sets `rg_reading` to False, disabling further reads and enabling the write process to run.

The next two rules, `rl_get_x1` and `rl_get_x2...` also alternate. The first one receives a response (from the first input vector), holds it in `rg_x1`, and hands off to the second rule, which receives another response (from the second input vector), performs the additions, and buffers the element sum in `fifo_sums`. It also accumulates the overall partial sum in `rg_partial_sum`.

```
———————————————————— App_HW_v0.bsv ————————————————————
  rule rl_get_x1 ((! rg_x1_valid) && (fifo_rsps.first.cmd == RSP_RD_DATA));
     rg_x1       <= fifo_rsps.first.data; fifo_rsps.deq;    // from vin1
     rg_x1_valid <= True;
  endrule


  rule rl_get_x2_sum_and_gen_wr_reqs (rg_x1_valid
                                        && (fifo_rsps.first.cmd == RSP_RD_DATA));
     let x1 = rg_x1;
     let x2 = fifo_rsps.first.data;  fifo_rsps.deq;    // from vin2
     let x3 = x1 + x2;

     rg_x1_valid <= False;

     // Write x3 to vout.
     fifo_sums.enq (x3);

     // Accumulate x3 in partial sum
     rg_partial_sum <= rg_partial_sum + x3;
  endrule
```

Rule `rg_wr_gen` executes a write burst, writing out the values in `fifo_sums` to memory. It uses `rg_wr_pc` to count of 1536 writes, after which it sets `rg_reading` to True, handing off control back to the read process.

The final rule, `rg_wr_rsp_drain` merely drains the right number of memory responses to write requests (thereby also assuring that all the writes have reached memory).

The remaining code defines the `start`, `result` and `mc_client` interfaces at the bottom of the module, which are all very straightforward. The `mc_client` interface simply connects to the relevant FIFOs.

3.2.3   Compiling and running the C and BSV code in Bluesim

First, please ensure that the Bluespec tools are installed and available on your development machine (which does not have to be the Convey HC server and does not need the Convey PDK). Please also ensure that you've defined the environment variables as usual, which will look something like this (modified as necessary for your Bluespec installation path and Bluespec release number):

```
———————————————————— Bluespec installation ————————————————————
   BLUESPEC_HOME=<path>/Bluespec-2014.07.A
   BLUESPECDIR=<path>/Bluespec-2014.07.A/lib
   BLUESPEC_LICENSE_FILE=@licenseServer.myInstitution.com
   PATH=....:/Bluespec-2014.07.A/bin:....
```

Details of getting to this point are given in the Bluespec User Guide in your installation. Also, please define the following environment variable to point at your BClib installation (modified as necessary for your BClib installation path and BClib release number):

```
─────────── Bluespec installation ───────────
    BCLIB=<path>/BClib-2014.12.C/MX        // for MX platforms
    BCLIB=<path>/BClib-2014.12.C/MX        // for HC platforms
```

Next, please copy the entire Example directory from your installation into a personal working directory where you have read/write privileges:

```
─────────── Make private copy of Example directory ───────────
[MyWorkArea] % cp -r  $BCLIB/Example  ./Example
```

In the BSV source directory, create or ensure that you have a symbolic link to version 0 of the example:

```
─────────── Properly named App_HW.bsv application source file ───────────
[Example       ] % cd BSV_src
[Example/BSV_src] % ln -f -s  App_HW_v0.bsv  App_HW.bsv
```

Finally, back in the Example directory, build the simulation executable:

```
─────────── Build Bluesim executable ───────────
[Example] % make  compile  link
```

Please feel free to examine the Makefile to see what this does. Briefly,

- It uses Bluespec's *bsc* tool to compile the $(BCLIB)/BClib_src/BC_Sim_Main.bsv file which, in turn, will cause the compilation of the App_HW.bsv file, and so on.
- It C-compiles the file C_src/App_SW.c using the "cc" compiler command.
- It uses Bluespec's *bsc* tool to link all the objects (from BSV and from C) into a standard Linux executable, called out (along with a shared object file called out.so).

You can now run this simulation executable:

```
─────────── Run Bluesim simulation ───────────
[Example] % make  simulate
```

or, just say:

```
─────────── Run Bluesim simulation ───────────
[Example] % ./out
```

We have provided a text file transcript_v0_bsim that shows what you ought to see as as terminal output from the simulation (it will likely differ in details like memory addresses).

This simulation of BSV code is 100% cycle-accurate with respect to the Verilog generated for deployment on the hardware. You can see the waveforms by re-simulating with the standard Bluesim -V flag:

```
─────────── Bluesim producing VCD waveforms ───────────
[Example] % ./out -V
```

which will dump a VCD file `dump.vcd` which can be viewed in your favorite waveform viewer (the -V flag optionally takes a preferred name for the VCD file).

Comment: The simulation just described does not attempt to model other aspects of the hardware: the memory system attached to the FPGAs (8 banks, pipelined crossbars, arbitration, DRAM access behavior, ...), Convey's PDK Verilog that surrounds the BSV design on the FPGA, the communication hardware between the x86 and the scalar co-processor and the FPGAs, the particular host x86 and memory system on the Convey HC server, etc. Thus, the overall system simulation in Bluesim will of course not be cycle-accurate with respect to the overall system execution on the Convey HC. Here, cycle-accuracy just refers to the cycle-equivalent behavior of the BSV code simulation and the corresponding generated Verilog, and the corresponding generated synthesized gates.

3.2.4 Compiling BSV code to Verilog, prior to deploying on Convey HC/MX hardware

Once you are satisfied with your application in simulation (previous section), you're ready to deploy it on the actual HC/MX hardware. Only Step 1, described here, is part of BClib. Until this step your development can be on any machine with Bluespec installed (i.e., it does not have to be on your Convey server, nor does it need the Convey PDK). Steps 2 onward are described separately in Appendix A, because those are standard Convey PDK steps and not part of BClib.

Step 1: Compile your BSV files to Verilog

This step is done using:

```
───────────────────────────── Creating Verilog ─────────────────────────────
[Example] % make verilog
```

It invokes *bsc* to compile, into Verilog, the BClib-supplied file `Cae_Pers.bsv` which, in turn, causes the compilation to Verilog of the `Example/BSV_src/App_HW.bsv` file, etc. The flags given to *bsc* cause the generated Verilog files to be placed in the directory `Example/Personality/verilog/`. One of the generated files is `cae_pers.v`, which is required by Convey's PDK as the user's top-level Verilog file.

After this, please follow Step 2 onwards in the Convey PDK standard flow, described in Appendix A.

3.3 Vector-add example, version 1

Like version 0, this version also assumes the start of the input and output vectors are aligned with MC 0. It also splits the work into per-MC engines, but for illustrative purposes it does it differently—it has 8 engines, each of which accesses both the even and odd ports of the MC.

Unlike version 0, we do not assume in-order memory responses and, instead, we instantiate a BSV memory-reordering transactor in the top module.

We will go over this version more quickly because the only difference is in the BSV HW implementation.

3.3.1 SW code walk-through

The SW side is identical to version 0; please refer to Sec. 3.2.1.

3.3.2   HW code walk-through

Please refer to the code which is in the file: `Example/BSV_src/App_HW_v1.bsv`
This is about 415 lines of BSV. We will just focus on the interesting differences from version 0. Like version 0, the file describes modules of a 3-level module hierarchy:

```
                                        App_HW_v1.bsv
module mkApp_HW (BC_HW_IFC);
   ...
endmodule
...
module mkApp_HW2 (BC_HW2_IFC);
   ...
endmodule
...
module mkVadder (Vadder_IFC);
   ...
endmodule
```

The top module `mkApp_HW` is similar to version 0, except that we instantiate a memory ordering transactor in a line like this:

```
                                        App_HW_v1.bsv
   match { .v16_mc_clients, .v16_mc_servers }
      <- mkBC_Vec_MC_Port_Ordering_Transactors;
```

Then, we connect the second-level module to the "server" side of these transactors:

```
                                        App_HW_v1.bsv
   BC_HW2_IFC hw2 <- mkApp_HW2;
   mkConnection (hw2.mc_ifcs, v16_mc_servers);
```

And, at the end of the module, we use the "client" side of these transactors to connect to the outside:

```
                                        App_HW_v1.bsv
   interface mc_ifcs        = v16_mc_clients;
```

The rest of `mkApp_HW` and `mkApp_HW2` are very similar to their counterparts in version 0. The `mkChanVadd` module has a slightly different structure from the counterpart in version 1. Instead of having 16 of these, each talking to an even or odd port of one of the 8 MCs, we have 8 of these, each talking to both the even and the odd port of one of the 8 MCs. Reads from input vectors 1 and 2 are launched in parallel into the even and odd ports, respectively. This is embodied in the two uses of `mkMC_Client` in the interface section of the module. The rule `rl_sum_and_gen_wr_reqs` receives both responses in parallel, performs the computation, and launches successive writes alternately into the even and odd ports of the MC.

3.3.3   Compiling and running the C and BSV code in Bluesim

Before you build for simulation or Verilog generation, please remember to change the symbolic link to point to version 1 of the BSV HW:

```
                    Properly named App_HW.bsv application source file
[Example        ] % cd BSV_src
[Example/BSV_src] % ln -f -s  App_HW_v1.bsv  App_HW.bsv
```

The actual build and run steps are identical to version 0; please refer to Sec. 3.2.3.

3.3.4 Compiling BSV code to Verilog, prior to deploying on Convey HC/MX hardware

Before you build for simulation or Verilog generation, please remember to change the symbolic link to point to version 1 of the BSV HW:

```
─────────────── Properly named App_HW.bsv application source file ───────────────
[Example        ] % cd BSV_src
[Example/BSV_src] % ln -f -s  App_HW_v1.bsv  App_HW.bsv
```

The actual build and run steps are identical to version 0; please refer to Sec. 3.2.4.

3.4 Vector-add example, version 2

In this version, unlike versions 0 and 1, we no longer assume that the input or output vectors have any particular alignment to the MCs or to each other. Thus, read and write requests and responses must be "routed" to the correct MC and back. We will go over this version more quickly because the only difference is in the BSV HW implementation.

3.4.1 SW code walk-through

The SW side is identical to version 1; please refer to Sec. 3.2.1.

3.4.2 HW code walk-through

Please refer to the code which is in the file: `Example/BSV_src/App_HW_v2.bsv`
This is about 500 lines of BSV. We will just focus on the interesting differences from the previous versions. Like those versions, the file describes modules of a 3-level module hierarchy:

```
─────────────────── App_HW_v2.bsv ───────────────────
module mkApp_HW (BC_HW_IFC);
   ...
endmodule
...
module mkApp_HW2 (BC_HW2_IFC);
   ...
endmodule
...
module mkVadder (Vadder_IFC);
   ...
endmodule
```

The top module `mkApp_HW` is identical to version 0. However, unlike version 0 where we assumed memory ordering (for which we relied on the Convey PDK facilities), here we do not assume any memory ordering. We will take care of memory ordering in (each instance of) `mkVadder`.

Later in the file, you will see lines like this:

```
─────────────────── App_HW_v2.bsv ───────────────────
// ----------------
// # of Vadders
typedef  8  M;    // # of VAdders
typedef  TAdd #(M,1)   M1;
```

In Version 0 (and 1), we had 16 (and 8) per-MC adders because it was so tightly structured around the raw memory architecture, which has 8 MCs. Here, since we are routing requests and responses anyway, we can independently choose how many adders we want (< 8 or > 8); this example chooses 8 (M). The second line defines M1 to be the type '9'.

In `mkApp_HW2` we instantiate an even and an odd "memory tree" root:

```
──────── App_HW_v2.bsv ────────
   MemTreeNode #(M, 8)    mem_tree_e <- mkMemTreeRoot_M_8;
   MemTreeNode #(M1, 8)   mem_tree_o <- mkMemTreeRoot_M1_8;
```

The modules `mkTreeRoot_M_8` and `mkTreeRoot_M1_8` are defined later in the file, but they're really just thin wrappers around the BClib module `mkTreeNode` (for separate synthesis) which are used to make "fat tree" memory routing networks. These modules are in the file `$BCLIB/BClib_src/BC_Mem_Tree.bsv`. The modules instantiated here have M (=8) root-side ports and M1 (=9) leaf-side ports. The root-side ports become the 8 MC client even/odd pairs. The M1 leaf-side ports connect to the M even and odd memory ports for the M adders, and the extra port is used by `mkApp_HW2` to read the param block, using the functions `send_rd_req`, `recv_rd_rsp` and `send_wr_req`, which follow next in the code.

Next, we instantiate the M adders, and connect their memory ports to the tree root ports:

```
──────── App_HW_v2.bsv ────────
   Vector #(M, Vadder_IFC) v_vadders <- replicateM (mkVadder);


   for (Integer m = 0; m < iM; m = m + 1) begin
      let mem_server_pair = tuple2 (mem_tree_e.leafside [m],
                                    mem_tree_o.leafside [m]);
      mkConnection (v_vadders [m].mem_client_pair, mem_server_pair);
   end
```

What follows is an FSM definition that reads the param block parameters, and then initiates the computations in M adders using their `start` method. Along the way, it computes a size that is roughly $\frac{1}{M}$, the segment of the vectors for which each adder is responsible:

```
──────── App_HW_v2.bsv ────────
          BC_Addr vsize_j = (rg_vsize / fromInteger (iM)) << 3;
```

Since the $\frac{1}{M}$ division may not be exact, the last adder is outside the for loop and is given responsibility for all the remaining elements. The rest of this module `mkApp_HW2` is straightforward and similar to the `mkApp_HW2` in version 1.

The `mkVadder` module that follows computes the partial sum for a segment of the vectors. The line:

```
──────── App_HW_v2.bsv ────────
   CompletionBuffer2 #(100, BC_Data) cb_rsps_1 <- mkCompletionBuffer2;
```

instantiates a module responsible for re-ordering memory responses that may arrive out of order. This module is found in the `$BCLIB/BClib_src/` directory (it is a small variation on the similar `mkCompletionBuffer` module in the standard Bluespec library and documentation). It is like a queue or FIFO in which we can "pre-reserve" an ordered spot so that when the data for that spot arrives, we can insert it in the correct place. The parameter 100 specifies how many reads may be in flight simultaneously (each such read has a reserved buffer, so a large parameter value can support more items in flight, but will cost more in hardware storage resources). We've chosen it to be 100 here because Convey FPGA memory latency is typically a 100 cycles or more.

The rule `rl_gen_rd_addrs` generates the read requests, launching them into the even and odd memory request FIFOs. Its work is done in the function `fn_gen_rd_req`.

```
─────────────────────────────── App_HW_v2.bsv ───────────────────────────────
   function ActionValue #(BC_MC_REQ)
       fn_gen_rd_req (Reg #(BC_Addr) rg_rd_base,
                      Reg #(BC_Addr) rg_rd_limit,
                      CompletionBuffer2 #(100, BC_Data) cb_rsps);
     actionvalue
        let cbtoken <- cb_rsps.reserve.get;
        let req  = BC_MC_REQ {cmd_sub: REQ_RD, rtnctl: extend (pack (cbtoken)), len:BC_8B,
                              vadr: rg_rd_base, data:?};
        let next_base = rg_rd_base + 8;
        rg_rd_base   <= next_base;
        return req;
     endactionvalue
   endfunction
```

It obtains a token to reserve an ordered spot in the return queue (`cb_rsps`), and sends this token out with the request. It increments the address by 8 (bytes) because here in version 2 we are no longer concerned with MC alignment, and this adder is responsible for a contiguous-address segment of this vector.

Later, in the `put` method, when the memory response arrives with the token, the token is used to place it into the correct spot in the return queue `cb_rsps`. The rule `rl_sum_and_gen_wr_reqs` drains the return queue in order, performs the additions, and launches the write request.

3.4.3   Compiling and running the C and BSV code in Bluesim

Before you build for simulation or Verilog generation, please remember to change the symbolic link to point to version 2 of the BSV HW:

```
──────────── Properly named App_HW.bsv application source file ────────────
[Example         ] % cd BSV_src
[Example/BSV_src] % ln -f -s  App_HW_v2.bsv  App_HW.bsv
```

The actual build and run steps are identical to the previous versions; please refer to Sec. 3.2.3.

3.4.4   Compiling BSV code to Verilog, prior to deploying on Convey HC/MX hardware

Before you build for simulation or Verilog generation, please remember to change the symbolic link to point to version 2 of the BSV HW:

```
──────────── Properly named App_HW.bsv application source file ────────────
[Example         ] % cd BSV_src
[Example/BSV_src] % ln -f -s  App_HW_v2.bsv  App_HW.bsv
```

The actual build and run steps are identical to the previous versions; please refer to Sec. 3.2.4.

A Deploying on Convey HC/MX hardware

> *Note: The following steps are standard Convey PDK flow steps. They are described here briefly for your convenience and for completeness of this manual, but the definitive reference is the Convey PDK Reference Manual. Convey's PDK offers many other build options, which are listed in their documentation.*
>
> *Any problems encountered in these steps are perhaps best resolved with Convey Support* (`support@conveycomputer.com`) *and not with Bluespec Support.*

Prior to this, you should have completed the following:

- Built and run your application in Bluesim (Sec. 3.2.3).
- Generated Verilog from your BSV code into the `Example/Personality/verilog/` directory, as described in Sec 3.2.4 "Step 1".

Step 2: Check that `Makefile.include` is ok for your site (one-time step)

Examine and possibly edit the file: `Example/Personality/Makefile.include`.

This file is adapted from the `Makefile.include` file provided by Convey in its PDK for its Verilog vector-add example. The first part is:

```
───────────────────── Makefile.include ─────────────────────
# ----------------------------------------------------------------
# This section added to original Convey Makefile.include, for BClib builds


# This allows Xilinx XST to exploit two cores (for faster synthesis)
export XILINX_THREADS = 2


# This points at Verilog files in Bluespec's standard library
USER_VERILOG_DIRS += $(BLUESPECDIR)/Verilog


# This tells XST to ignore main.v in Bluespec's standard library
USER_VERILOG_SKIP_LIST += $(BLUESPECDIR)/Verilog/main.v


# This points at the BClib library
USER_VERILOG_DIRS += $(BCLIB)/BClib_verilog
# ----------------------------------------------------------------
```

Ensure the next two definitions are correct for your installation (or edit them as necessary):

```
───────────────────── Makefile.include ─────────────────────
    CNY_PDK     = /opt/convey/pdk_mx
    CNY_PDK_REV = latest
```

`CNY_PDK_PLATFORM` should be set to one of hc-1, hc-1ex, hc-2, hc-2ex, or MX.

`PERS_IF_TYPE` should be set to 0 for HC platforms and 1 for MX platforms.

`MC_XBAR` refers to Convey's optional on-FPGA routing logic for memory ports. When set to 0 (omitting the crossbar logic), requests for an address with address-bits $[8:6] = j$ should only be sent to ports $2j$ or

$2j + 1$ (out of 16 ports). When set to 1 (including the crossbar logic), any request can be sent to any of the 16 ports.

For all three versions of the BClib examples set `MC_XBAR = 0`.

`MC_READ_ORDER` refers to Convey's optional logic for ensuring that memory responses are returned in the same order as original requests. For BClib example version 0, we build with `MC_READ_ORDER=1`, (include the logic). For BClib examples version 1 and 2, we build with `MC_READ_ORDER=0` (exclude the logic).

The whole `Personality` directory follows the structure recommended in Convey's PDK for user projects, and is in fact derived from the standard Convey PDK project example.

Step 3: Create the FPGA bitfile ("the long pole in the tent")

CAUTION: this step can take an hour or more as it involves Xilinx synthesis!
(It is the only long step in the whole flow.)

Ensure that your Convey PDK installation is ready, and that your Xilinx tools installation is ready (since the Convey PDK commands invoke Xilinx tools). Xilinx tools may need the `LM_LICENSE_FILE` environment variable to be set properly.

We recommend you only use XST version 14.2 or later, since earlier versions have a bug in the Verilog parser that can result in wrong synthesized circuits for the FPGA.

Start with a clean `Personality/phys` directory:

```
───────────────── Clean the phys directory ─────────────────
[Example/Personality/phys] % make clean
```

Extra step only for users of Virtex-5 based Convey platforms, such as HC-1 or HC-2

Even with Xilinx XST 14.2, when targeting a Virtex-5 device, it still uses XST's older, buggy, Verilog parser by default. Please force it to use its new Verilog parser as follows. Do the following "make" command which will create the file `cae_fpga.xst` in the `phys` directory:

```
[Example/Personality/phys] % make cae_fpga.xst
```

Edit the `cae_fpga.xst` file, adding the following line:

```
-use_new_parser YES
```

(If you have an HC-1ex, HC-2ex or MX, you can ignore the step in this box, since XST 14.2 uses the new parser by default for Virtex-6).

Finally, a bare "make" command invokes the PDK procedures to create the bitfile:

```
───────────────── Creating the FPGA bitfile ─────────────────
[Example/Personality/phys] % make
```

This invokes the Xilinx synthesis tool and creates a bitfile inside the `phys` directory. You should always check the synthesis reports to ensure that the bitfile is able to run at 150 MHz, the default timing target set by the PDK tools (see box on "Timing Closure").

Timing Closure

The Xilinx synthesis tool XST is invoked when you do a 'make' in the `Personality/phys` directory to create your bitfile. It also places all Xilinx synthesis reports in this directory. You should examine the timing report file `cae_fpga_routed.twr` (Warning: this is often a very large text file), for lines similar to this:

Slack: -0.2ns (requirement - (data path - clock path skew + uncertainty))

A negative slack like this means your design did not meet timing, i.e., there is some circuit path that overshot the default 6.667 nanosecond target (= 150 MHz) by, in this case, 0.2ns. There may be dozens, or hundreds of lines like this, but it is usually sorted from worst negative slack downwards. If you have any such paths, you should fix them before proceeding.

For experimental runs during development, your bitfile may still run if you have paths with negative slack up to about 10% of the target period, i.e., up to about -0.667ns. For production releases you should aim for *all* paths having positive slack.

To fix timing problems:
- Each negative slack report is immediately followed by a listing of the offending critical path.
- From the start and end points of the critical path, you will be able to identify the corresponding BSV source code state elements (these names are preserved from BSV into Verilog, and thus into the Xilinx critical path description). This will give you an indication of what needs to be changed in the BSV source to shorten the path. The intermediate gates named on the critical path may also give you a sharper indication of what needs to be changed.
- Fix your BSV source code appropriately (use smaller expressions, move some logic into another rule, split a pipeline stage into two, ...). You should probably Bluesim-simulate again, to ensure it is still functionally correct, and then recreate the Verilog and make the bitfile again. *We do not recommend* ever *editing the bsc-produced Verilog files; always fix it in the BSV source!*
- "Repeat and Rinse" until your synthesis meets timing.

Note: the default 150 MHz timing target is based on the speed of the memory ports in the Convey hardware. Both versions of the BClib examples meet this target. If meeting this target is too difficult for your design, Convey's PDK offers an option to run your hardware design at a lower clock speed, and to insert appropriate clock-change logic in the interfaces to the memory ports. Please see Convey PDK documentation for more information, or ask Convey Support for help.

Note: the Xilinx reports also include a file called `cae_fpga.srp` which, among other things, describes the timing achieved for your design. Please *do not* rely on this—it is only a very rough estimate made in the early stages of Xilinx synthesis. It can be as much as 100% off. Use the timing reports in the `.twr` file only.

Step 4: Create the Convey PDK bitfile release

Do the following Convey PDK command:

```
──────────── Creating the PDK bitfile release ────────────
[Example/Personality/phys] % make release
```

This will create a timestamped release:

```
──────────── The timestamped bitfile release ────────────
    Example/Personality.released/12_08_14_59/
        ae_fpga.tgz
        Personality_12_08_14_59.tar.gz
        rpm/
```

The directory name (yy_hh_mm_ui) represents the date on which you did your build, suffixed with a unique id that is regenerated each time you do a build. You will be copying `12_08_14_59/ae_fpga.tgz` (directory with one file) to your Convey server in a later step.

Step 5: Choose a personality number and nickname and create a personality directory (one-time step)

Pick a personality number and nickname for your custom personality, such as 65123 and "foo". Convey suggests using numbers $\geq$ 65000 for site-local and experimental personalities[1]. On the Convey HC server, create a directory for your custom personality, such as:

```
─────────────────── Personality directory ───────────────────
[HOME] % mkdir  personalities/65125.1.1.3.0
```

The '3' indicates an MX machine. For HC-1 and HC-2 machines, use 65125.1.1.1.0, i.e., change the last '3' to a '1'. makeFor HC-1ex and HC-2ex machines, use 65125.1.1.2.0, i.e., change the last '3' to a '2'.

Copy the contents of Convey's example personality (4.1.1.1.0 or 4.1.1.2.0) into this directory, and create a "customdb" file listing your new personality number and nickname:

```
─────────────── Personality directory and customdb file ───────────────
[HOME] % cp -r  /opt/convey/personalities/4.1.1.1.0/*   ./personalities/65125.1.1.3.0/
[HOME] % echo "65125.1.1.3.0,foo" > personalities/customdb
```

Set the environment variable that tells the Convey tools and run-time systems where to look for your personality:

```
─────────────── Personality environment variable ───────────────
[HOME] % export CNY_PERSONALITY_PATH=$HOME/personalities
```

Step 6: Install your release bitfile on Convey HC/MX server

Copy your timestamped bitfile release into the personality directory

```
─────────────── Personality bit file location ───────────────
    $HOME/personalities/65125.1.1.3.0/12_08_14_59/ae_fpga.tgz
```

and make a symbolic link to it from the file `ae_fpga.tgz`:

```
─────────────── Installing each timestamped personality bit file ───────────────
[HOME] % cd  personalities/65125.1.1.3.0/
[HOME/personalities/65125.1.1.3.0] % ln -s -f 12_08_14_59/ae_fpga.tgz  ae_fpga.tgz
```

As you debug and evolve your application and create new timestamped releases, you only have to repeat this step to install it.

---

[1] If you expect more widespread deployment of your application to multiple Convey servers, customer Convey servers, remote Convey servers, etc., the process is more involved. You will have to obtain a unique personality number from Convey, and there are more packaging steps. Please consult Convey PDK documentation and Convey Support.

Step 7: Compiling and linking your C application program

You will need your C application files (`App_SW.c` for our examples), and copies of some BClib C and assembly language files (which are found in `$BCLIB/BC_src/`):

```
    cpAsm.s
    BC_linkage.h          BC_linkage.c
    timing.h              timing.c             // optional
    instrumentation.h     instrumentation.c    // optional
```

In each of the files `App_SW.c` and `cpAsm.s` you will find one line that mentions the personality number. Edit these, if necessary, to match your chosen personality number (e.g., 65125). This is how the Convey runtime system knows to associate this SW code with that particular HW implementation.

In any convenient working directory on the Convey HC machine, compile and link these files, as illustrated in the following sample Makefile entry:

```
───────────── Compiling and linking your SW program ─────────────
# Sample Makefile entry
UserApp.exe: App_SW.c  BC_linkage.c  BC_linkage.h  cpAsm.s
        cnycc  -DFOR_HW  App_SW.c  BC_linkage.c  cpAsm.s  -lpthread  -o UserApp.exe
```

It creates the executable program `UserApp.exe` by invoking the Convey C compiler `cnycc`. Note that we define the variable `FOR_HW` which is used in `BC_linkage.c` to invoke various Convey runtime library functions. It may also be used by the application program (in our example, `App_HW.c`, the classical `main()` entry point is ifdef'd based on this).

Step 8: Running the SW-HW application on the Convey HC/MX hardware

To run the program, we do:

```
──────────── Running the SW-HW application on HC/MX hardware ────────────
[WorkDir] % /opt/convey/sbin/mpcache -f
[WorkDir] % ./UserApp.exe
```

The first is a Convey command to clear the bitfile cache in the HC. It only needs to be done if you have an older version of the bitfile with the same personality number loaded on the machine from a previous run.

Then, we invoke the SW executable. When the C code executes the BClib function `bc_init_HW ()` it will, via the Convey runtime library, identify the signature of the desired bitfile, which will get loaded on the FPGAs in the HW. When the C code executes the BClib function `bc_call_HW ()` it will, via the Convey runtime library and assembly code in `cpAsm.s`, communicate with the FPGA. When the FPGA code signals completion, it returns from the function, and then the C code completes and returns you to the command-line prompt.

Note: bitfile loading only happens the first time you run an app with a new signature. When the bitfile for that signature is already loaded, subsequent runs just use it.

B Compiling BSV code to Verilog and running the app on Convey's PDK Simulator

[*This is optional, and described here only for completeness. We typically only do Bluesim simulation as described in Sec. 3.2.3, and then build and deploy on the Convey HC/MX hardware. Bluesim's cycle-accurate simulation of your application BSV code essentially plays the same role as Convey's PDK Verilog-based simulation (and typically at much higher simulation speeds).*]

If you wish to run your application in the Convey PDK simulation environment, you must also have one of the Convey-recommended Verilog simulators such as Modelsim or VCS installed, along with the proper license for the simulator.

You will just need the Verilog generated from `App_HW.bsv` (as described in Sec. 3.2.4), and the C and assembly code files described in Appendix A Step 7. Please refer to Convey PDK documentation for details, but here is a brief description of setting up a PDK sim. Copy the `SampleAppVadd` and `sim` directories from the Convey PDK project example into your `Personality` directory:

```
─────────────────── Convey PDK sim setup ───────────────────
% cp -r /opt/convey/pdk/version/<platform>/examples/cae_vadd/sim            .
% cp -r /opt/convey/pdk/version/<platform>/examples/cae_vadd/SampleAppVadd  .
```

(with `<platform>` as usual being hc-1 or hc-1ex.) Edit the file `Personality/sim/sc.config`. Look for the line:

```
─────────────── Edit this line in the file sc.config ───────────────
    caesim  2  0
```

Change the '0' to '1'. (With '0', the PDK simulator expects also another software program that functionally models your hardware, which it runs in parallel with your Verilog simulation, and it compares them. This edit disables this feature.).

In the `SampleAppVadd` directory, do the actions of Appendix A Step 7 to build the executable: copy files `App_SW.c`, `BC_linkage.h`, `BC_linkage.c` and `cpAsm.s`, edit `App_SW.c` and `cpAsm.s` if necessary to have your chosen personality number, compile and link them into an executable `UserApp.exe`. (And even though you will be running in Verilog simulation, you still need the `-DFOR_HW` compilation flag described in Step 7).

Finally, run your SW executable `UserApp.exe` using the run script in the example:

```
─────────────────── Convey PDK sim run ───────────────────
[Personality/SampleAppVadd] % ./run
```

C Some performance tips

When measuring performance of your app on the hardware, please be aware of the following. None of these are particular to BClib; we have observed the same behavior with Convey's PDK vector-add example.

- When you run an app requiring a particulare signature, the bitfile is loaded by the Convey run time system only if it is not already loaded on the hardware from the previous run. Thus, there is a "first time" performance penalty, which can be significant. We recommend doing performance measurements only on a second or later run when the bitfile is already loaded.

- The data structures accessed by the FPGAs should be allocated on the FPGA-side of the machine (see `bc_malloc` in the BClib example C code). There is a significant performance penalty in accessing "the other side" (C code on x86 accessing FPGA-side memory, or FPGA accessing x86-side memory).

- Even when your data structures are allocated on the FPGA side, there seems to be a "first touch" penalty until all the data has actually been accessed by the FPGA. In the vector-add examples (both in BClib and in Convey's PDK example), if we surround the main computation in a for-loop that runs it 10 times, we observe significantly larger run times for the first run compared to the remaining 9. We are not sure of the reason for this; perhaps the Convey run-time system really maps the pages of the allocated memory only on demand?

- Each of the 16 memory ports on an FPGA is used for both read and write requests. However, there seems to be a performance penalty if a read is immediately followed by a write or vice versa (perhaps the direction of some internal bus in the Convey PDK has to be turned). Thus, it is advantageous to have contiguous sequences of reads and contiguous sequences of writes. In the `App_HW_v0.bsv` example, we do 1536 consecutive pairs of reads followed by 1536 consecutive writes; this produced a 40% improvement in application performance.

BClib Reference Manual

D   Reference Manual Overview

Sec. D.1 describes the basic BSV interfaces and related types.

Sec. D.2 describes some utility functions that are useful for all apps.

Sec. D.3 describe some convenient, pre-defined transactors that further simplify application code.

For building and executing using Bluesim, please refer to the detailed example in Sec. 3.2.3.

For building and executing on actual Convey HW, please refer to the detaile example in Sec. 3.2.4.

Appendix B summarizes how to build and execute in the Convey PDK simulation environment (using Verilog simulation).

All the BClib source files described below are in the BClib installation directory, at `$BCLIB/BClib_src/`. A user of BClib is expected to examine only the file `BC_HW_IFC` in detail; the rest of the sources are provided mainly for reference, and for users who choose to write their own variations.

D.1   Top-level of the app's BSV code and its interface

The top-level of the app's BSV code must be a module named `mk_App_HW`. It must implement the interface `BC_HW_IFC`, which is defined in the file `BC_HW_IFC.bsv`, and which is typically imported into the app's BSV code using:

```
import    BC_HW_IFC :: *;
```

General note: various types below have an `instance FShow` declaration. This indicates that the `fshow()` function can be used to "pretty-print" values of that type in `$display` statements.

D.1.1   Single and multiple FPGAs

The Convey hardware has four user-programmable FPGAs (a.k.a. AEs or Application Engines). The descriptions below are for a single FPGA. You can use them to create a single HW design that is replicated across the four FPGAs, or create separate four separate designs, one for each FPGA. This choice is outside the scope of discussion of this reference manual; please see the Convey PDK Reference Manual for a discussion on how to load the four FPGAs with replicated or differentiated designs.

D.1.2   Interface `BC_HW_IFC`

This is the top-level interface that should be implemented by the application's BSV code:

```
──────────────────── BC_HW_IFC.bsv ────────────────────
interface BC_HW_IFC;
   // Memory interfaces
   interface Vector #(16, BC_MC_Client)  mc_ifcs;
```

```
    interface BC_Dispatch_IFC    dispatch_ifc;
    interface BC_Management_IFC  management_ifc;
endinterface: BC_HW_IFC
```

These correspond directly to the three interface groups provided by the Verilog infrastructure in Convey's PDK (Dispatch, MC and Management), except that here they are presented as BSV "transactional", flow-controlled interfaces rather than raw Verilog signals. Note that the MC sub-interface is a vector of 16 `BC_MC_Client` interfaces, which connect to the underlying Convey PDK Verilog MC interfaces (8 x even/odd). These provide full access the Convey PDK facilities.

BClib does not provide access to the direct AE-to-AE communication links. AE-to-AE communication can always be done through memory.

D.1.3 The Dispatch sub-interface `Convey_Dispatch_IFC`

```
———————————————————— BC_HW_IFC.bsv ————————————————————
interface BC_Dispatch_IFC;
    interface Put #(BC_Dispatch_inst)      put_inst;
    interface Get #(BC_Dispatch_ret_data)  get_ret_data;
    interface Get #(BC_ExceptionVec)       get_exception;


    (* always_ready *)
    method    Bool                         idle;
    (* always_ready *)
    method    BC_Aeg_Cnt_Rep               aeg_cnt;
```

It has sub-interfaces using which the environment can insert ("put") instructions, and retrieve ("get") return-data and exception vectors. These are fully flow-controlled, i.e., the user does not have to worry about handshake signals to reliably transfer instructions or return-data.

The `idle` method tells the environment when the module is idle (no instructions currently in flight), and the `aeg_cnt` method tells the environment the number of AE General-Purpose registers implemented by the module. The types of data returned for aeg counts and exceptions are:

```
———————————————————— BC_HW_IFC.bsv ————————————————————
typedef  Bit #(18)  BC_Aeg_Cnt_Rep;
typedef  Bit #(16)  BC_ExceptionVec;


// Standard exceptions


// Bit 0, aka AEUIE
BC_ExceptionVec bc_exception_unimplemented_inst = 16'h0001;
// Bit 1, aka AEERE (Element Range Exception)
BC_ExceptionVec bc_exception_invalid_aeg_idx    = 16'h0002;
```

Exception vectors are 16-bit vectors. Convey pre-defines the meaning of the bottom two bits (unimplemented instructions, and invalid AEG index). The remaining 14 bits are application-specfic.

Instructions inserted by the environment, and return-data sent to the environment, are defined by the following types:

```
——————————————————————————————— BC_HW_IFC.bsv ———————————————————————————————
typedef  Bit #(32)  BC_Inst;
typedef  Bit #(64)  BC_Data;

// Dispatch instructions, as seen by BSV app
typedef struct {
   BC_Inst  inst;
   BC_Data  data;
   } BC_Dispatch_inst
   deriving (Bits);

instance FShow #(BC_Dispatch_inst);

// Dispatch ret_data, as seen by BSV app
typedef BC_Data  BC_Dispatch_ret_data;
```

D.1.4   The MC sub-interface `BC_MC_Client`

Memory Controller Ports, Addressing and Routing

In Sec. D.1.2 we saw that the interface to memory consists of a vector of 16 "clients" which connect to the 8 MCs. Each MC has a pair of ports. The first 2 clients connect to the first MC's pair of ports, the second 2 clients connect to the second MC's pair of clients, and so on.

Memory is striped across 8 MCs either in a simple binary interleave, or in a so-called "31_31" interleave (see Convey PDK Reference Manual for a detailed discussion). Each memory request should be directed to the MC that implements the address in the request. For a simple binary interleave, bits [8:6] of the address identify which of the 8 MCs should be targeted (see also Sec. D.2.1 for some utility functions to assist in memory addressing and routing).

[Note: Convey also provides so-called "xbar" or cross-bar support, using which it is not necessary to direct a request to a particular MC, i.e., the xbar routes it to the right MC. Although more convenient to use, (a) this will cost FPGA area and (b) performance-critical designs will anyway want to understand MC routing to avoid unnecessary collisions. Hence the facilities in BClib do not assume any xbar support.]

Each of the 16 MC clients is defined as follows:

```
——————————————————————————————— BC_HW_IFC.bsv ———————————————————————————————
interface BC_MC_Client;
   interface Client #(BC_MC_REQ, BC_MC_RSP)            req_rsp;
   interface Client #(BC_MC_flush_req, BC_MC_flush_rsp)  flush;
endinterface
```

It contains sub-interfaces where the environment can get a memory request from, and put a read response into the module. The environment can also get a request to flush outstanding writes, and can put a response that the such a flush is complete. Everything is pipelined, i.e., each of these interfaces is a stream of requests or responses.

Memory requests and responses

Memory requests have the following type:

```
————————————————————————— BC_HW_IFC.bsv ——————————————————
typedef struct {
   BC_MC_REQ_TYPE  cmd_sub;     // 3-bit cmd and 4-bit sub
   BC_DataSize     len;
   BC_Addr         vadr;
   BC_RtnCtl       rtnctl;
   BC_Data         data;
   } BC_MC_REQ
   deriving (Bits, FShow);
```

The `cmd_sub` field uses one of the following labels corresponding to all the availble Convey memory "cmd"
and "sub" codes (see Convey MX reference manual):

```
————————————————————————— BC_HW_IFC.bsv ——————————————————
typedef enum {
   // Stores
   REQ_WR,
   // Loads
   REQ_RD,
   // Full Empty requests
   REQ_RFE, REQ_RFF, REQ_IFF, REQ_REF, REQ_WEF,
   REQ_WFF, REQ_WXF, REQ_WXE, REQ_RXX, REQ_PUR,
   // Atomic requests
   REQ_ATOM_ADD, REQ_ATOM_SUB, REQ_ATOM_EXCH, REQ_ATOM_MIN,
   REQ_ATOM_MAX, REQ_ATOM_INC, REQ_ATOM_DEC,  REQ_ATOM_AND,
   REQ_ATOM_OR,  REQ_ATOM_XOR, REQ_ATOM_CAS
   } BC_MC_REQ_TYPE
deriving (Bits, Eq, FShow);
```

The `len` field identifies whether requests are for 1, 2, 4 or 8 bytes:

```
————————————————————————— BC_HW_IFC.bsv ——————————————————
typedef  enum { BC_1B, BC_2B, BC_4B, BC_8B } BC_DataSize
   deriving (Eq, Bits, FShow);
```

The address, "return control" and data fields are defined as follows:

```
————————————————————————— BC_HW_IFC.bsv ——————————————————
typedef Bit #(48)  BC_Addr;
typedef Bit #(64)  BC_Data;


typedef 32 BC_RTNCTL_WIDTH;
typedef Bit #(BC_RTNCTL_WIDTH)  BC_RtnCtl;
```

The return control field is typically used because the path to memory and back is deeply pipelined. There
can be multiple requests in flight, and their responses may return in a different order. The user can
use the return control field to identify individual requests (e.g., by carrying a serial number), because it
accompanies the response. The number of requests in flight is implementation-dependent (on the Convey
HC-1, the number in flight can be as high as 256).

Memory responses are defined as follows:

```
                              ─────── BC_HW_IFC.bsv ───────
typedef struct {
   BC_MC_RSP_TYPE   cmd;
   BC_MC_RSP_SUB    sub;
   BC_RtnCtl        rtnctl;
   BC_Data          data;
   } BC_MC_RSP
   deriving (Bits, FShow);
```

The `cmd` field is the response op code and the `sub` field is the 3-bit additional information (see Convey MX reference manual):

```
                              ─────── BC_HW_IFC.bsv ───────
typedef enum {
   RSP_RD_DATA,   RSP_WR_CMP,      // load/store responses
   RSP_FE_RETRY,  RSP_FE_DATA,     // full/empty bit op responses
   RSP_ATOMIC_DATA                 // atomic op responses
   } BC_MC_RSP_TYPE
deriving (Bits, Eq, FShow);


typedef Bit #(3) BC_MC_RSP_SUB;
```

Flushes

Since the path to memory is deeply pipelined, the fact that your module has issued a write request is not an indication that the request has actually reached memory or that the data has been written. If it is important to be sure that all outstanding writes have completed, your module can issue a flush request and wait for a flush-complete response. You issue a flush request merely by returning a request on the `get_flush_req` interface to the environment (in the `BC_MC_Client` interface described earlier). A completion response arrives when the environment calls the `put_flush_rsp` method. Flush requests and responses are defined as follows:

```
                              ─────── BC_HW_IFC.bsv ───────
typedef struct {
   Bit #(1) dummy;
   } BC_MC_flush_req
deriving (Bits);

typedef struct {
   Bit #(1) dummy;
   } BC_MC_flush_rsp
deriving (Bits);
```

The actual 1-bit values don't really matter, as these are pure signals. To create a value of this type (e.g., when returning a flush request), use the "don't care" literal "?". They are wrapped in struct definitions so that they are not accidentally confused with other 1-bit values.

Note: In the original HC-1/HC-2 platforms, memory-writes had no responses, and flushes were the only way to guarantee that writes had completed. In the newer MX platform, memory-writes have responses, which guarantee completion. Thus, flushes seem to be of less use in the MX platform. None of the examples in BClib use flushes.

D.1.5   The Management sub-interface `BC_Management_IFC`

Caveat: This interface has only been partially defined in the current release. Specifically, the "CSR Ring" functionality is not implemented.

```
───────────────────────── BC_HW_IFC.bsv ─────────────────────────
typedef Bit #(2) BC_AEId;

interface BC_Management_IFC;
   (* always_ready, always_enabled *)
   method Action set_aeid (BC_AEId aeid);

   (* always_ready, always_enabled *)
   method Action m_csr_31_31_intlv_dis (Bool csr_31_31_intlv_dis);

   // ----------------
   // CSR ring
   (* always_ready, always_enabled *)
   interface Put #(BC_Ring_msg) put_ring_msg_in;

   (* always_ready, always_enabled *)
   interface Get #(BC_Ring_msg) get_ring_msg_out;
endinterface

typedef Bit #(4)  BC_RingCtl;
typedef Bit #(16) BC_RingData;

typedef struct {
   BC_RingCtl   ctl;
   BC_RingData  data;
   } BC_Ring_msg
   deriving (Bits);
```

The 2-bit `set_aeid` method is how the environment informs an AE (FPGA) whether it is number 0, 1, 2 or 3 of the four FPGAs. Thus, even though all four FPGAs may be loaded with the identical bitfile, they can have different behaviors because they can test this value, which will be different for each of them.

The `m_csr_31_31_intlv_dis` method is how the environment informs the AE whether the 31-31 interleave is disabled (therefore, simple binary interleave), or enabled.

The CSR Ring is not implemented yet. The "null" management transactor described in Sec. D.3.3 passes the aeid and 31_31 interleave status through, and simply "loops back" the CSR ring.

D.2   Utility functions

These utility functions are also in the file `BC_HW_IFC.bsv`. In this section we describe some functions that are of general use to designers of all apps for the Convey AE (FPGA).

The following definitions are useful to "stub" out interfaces that are unused (e.g., the flush and flush-complete-response interfaces, if you don't need them).

```
───────────────────────── BC_HW_IFC.bsv ─────────────────────────
Get #(t) getstub;
Put #(t) putstub;
```

```
Client #(t1,t2) clientstub;
Server #(t1,t2) serverstub;
```

The get stub never returns a value, the put stub discards all values.

D.2.1 Binary interleave address utility functions

There are 8 MCs (Memory Controllers), and each AE (FPGA) has 2 ports (even and odd) to each of them. Thus, each AE has 16 ports in total, all of which can be used in parallel.

Convey HC memory is byte-addressed. "Words", "DoubleWords" and "QuadWords" are 2, 4 and 8 bytes, respectively. Thus, QuadWord-aligned addresses have 3'b000 for the bottom three bits.


Figure 5: Addressing with the simple binary interleave

When the hardware is configured for the simple binary memory interleave, Fig. 5 shows how data is striped across MCs. Thus,

- Bits [8:6] of an address identify the MC in which it lies.
- Zeroing out bits [5:0] of an address will align it to the start of the 64-byte chunk in an MC.
- Zeroing out bits [8:0] of an address will align it to the start of the 512-byte stripe across the 8 MCs.

With this in mind, the following definitions should be clear:

```
──────────────────────── BC_HW_IFC.bsv ────────────────────────
typedef Bit #(3)  BC_MC;    // 0..7 (index for 8 memory channels)

function BC_MC bc_mc_of_addr (BC_Addr addr) = addr [8:6];

function Bool bc_is_aligned_to_1B               (BC_Addr addr) = True;
function Bool bc_is_aligned_to_2B               (BC_Addr addr) = (addr [0] == 0);
function Bool bc_is_aligned_to_4B               (BC_Addr addr) = (addr [1:0] == 0);
function Bool bc_is_aligned_to_8B               (BC_Addr addr) = (addr [2:0] == 0);
function Bool bc_is_aligned_to_stripe_MC_chunk (BC_Addr addr) = (addr [5:0] == 0);
function Bool bc_is_aligned_to_stripe           (BC_Addr addr) = (addr [8:0] == 0);


function BC_Addr  bc_offset_in_stripe_MC_chunk (BC_Addr addr) = extend (addr [5:0]);
function BC_Addr  bc_offset_in_stripe           (BC_Addr addr) = extend (addr [8:0]);
```

Suppose we have an array in memory starting at `vec_base_addr`. This function computes the smallest address ≥ `vec_base_addr` that lies in MC J:

```
──────────────────────── BC_HW_IFC.bsv ────────────────────────
function BC_Addr bc_base_addr_for_chan (Integer chanJ, BC_Addr vec_base_addr);
   BC_MC  chanJ_b3 = fromInteger (chanJ);
   BC_MC  chan_of_base = bc_mc_of_addr (vec_base_addr);
```

```
   Bit #(39) addr_bits_above_chan = vec_base_addr [47:9];
   return (  (chanJ_b3 < chan_of_base) ? { (addr_bits_above_chan+1), chanJ_b3, 6'h0}
         : (  (chanJ_b3 == chan_of_base) ?vec_base_addr
           :                                      {addr_bits_above_chan, chanJ_b3, 6'h0}));
endfunction
```

Suppose we have a QuadWord address `qw_addr`. This function returns the address of the next QuadWord in the same channel (MC):

———————————————— BC_HW_IFC.bsv ————————————————
```
function BC_Addr bc_next_addr_8B_in_chan (BC_Addr addr_8B);
   // h38 is offset of last quadword in 64B block in chan
   return (  (addr_8B [5:0] != 6'h38)
         ? (addr_8B + 8)
           // skip to next chunk in this channel
         : ( { (addr_8B + 512)[47:6], 6'h0 } ) );
endfunction
```

D.3 Convenience Transactors

The following code is provided in the file `BC_Transactors.bsv`. The declarations are available to user code by importing it:

```
import    BC_Transactors :: *;
```

This section describes library elements that are likely to be useful for most developers, encapsulating common functionality. Fig. 6 explains the concept. Each library element is a transactor—a module—that


Figure 6: Convenience Transactors (library)

the user can instantiate inside the top-level `mkHW_App` module. Each transactor provides two interfaces, one of which can directly be used as the corresponding required sub-interface of `mkHW_App` (we also refer to this as the "outward-looking" interface). The other sub-interface is available inside the module to be accessed by user code (we also refer to this as the "inward-looking" interface). For example, the Simple Dispatch transactor (more detail in Sec. D.3.1) implements AEG registers, instruction decode, execution of AEG read/write instructions, etc., leaving only the actual CAEP0/1 (co-processor) instructions to be dealt with by user-code.

The user can pick and choose which of these transactors are useful for his app (including none); they are just independent, optional library elements for convenience. The vector-add example in Sec. 3 uses all of them.

D.3.1   Simple Dispatch transactor

The Simple Dispatch transactor takes care of routine functions like implementing the AEG registers, decoding and executing AEG read/write instructions, etc. What remains for the user to implement is just the actual CAEP0/1 (co-processor) instructions. The module is declared like this:

```
─────────────────────── BC_Transactors.bsv ───────────────────────
module mkSimple_Dispatch  #(parameter BC_Aeg_Idx aeg_idx_hi)
                          (Tuple2 #(BC_Dispatch_IFC,
                                    Simple_Dispatch_IFC));
```

The parameter `aeg_idx_hi` is the highest index of the AEG register file), and thus specifies the number of desired AEG registers. Note: this transactor implements all the registers between 0 and `aeg_idx_hi`, even if the user-code uses them sparsely.

The outward-looking interface is the standard `BC_Dispatch_IFC`. The inward-looking interface is declared like this:

```
─────────────────────── BC_Transactors.bsv ───────────────────────
interface Simple_Dispatch_IFC;
   interface Get #(BC_Dispatch_inst)  get_caep_inst;
   interface Put #(BC_ExceptionVec)   put_caep_result;

   method BC_Data   rd_aeg (BC_Aeg_Idx  idx);
   method Action    wr_aeg (BC_Aeg_Idx  idx, BC_Data data);

   (* always_ready, always_enabled *)
   method Action    set_aeid (BC_AEId aeid);
endinterface: Simple_Dispatch_IFC
```

The first sub-interface is for user-code to get CAEP0/1 instructions (these will be the only instructions coming through this interface, since all other instructions are handled inside the transactor). User-code should implement the chosen semantics of those instructions, and signal completion by using the second sub-interface to return an exception vector (which will be all zeros if there was no exception). User code can access the AEG registers using the `rd_aeg` and `wr_aeg` methods. The `set_aeid` method is used by the transactor only for $display statements, and can be left unconnected (see management null interface in D.3.3 for where to obtain this value from).

The transactor raises the standard "unimplemented instruction" and "invalid aeg index" exceptions.

The transactor stalls instruction decode and execution when any CAEP0/1 is in flight in the user code, i.e., once it has delivered a CAEP0/1 instruction to user code, it stalls until the exception-vector response is returned. Thus, user-code can assume stable values for the AEG registers.

The transactor will not assert the `idle` signal to the environment if any CAEP0/1 instruction is still in flight.

D.3.2   MC (Memory Controller) null and ordering transactors

In the memory transactors, each inward-looking port interface, shown below, is simply the get/put "dual" of `BC_MC_Client` shown above in Sec. D.1.4:

```
────────────────────── BC_Transactors.bsv ──────────────────────
interface BC_MC_Server;
   interface Server #(BC_MC_REQ, BC_MC_RSP)              req_rsp;
   interface Server #(BC_MC_flush_req, BC_MC_flush_rsp)  flush;
endinterface
```

BClib provides two transactors: null and ordering. The null transactor is a simple "pass-through" (thus out-of-order responses remain out-of-order). The transactor offers the full complement of 16 ports:

```
────────────────────── BC_Transactors.bsv ──────────────────────
module mkBC_Vec_MC_Port_Null_Transactors
   (Tuple2 #(Vector #(16, BC_MC_Client),
             Vector #(16, BC_MC_Server)));
```

The servers are inward-looking to the app, and clients implement the outward-looking interfaces.

The ordering transactor alleviates the user code from having to worry about out-of-order responses, i.e., it returns responses in the same order as their corresponding requests:

```
────────────────────── BC_Transactors.bsv ──────────────────────
module mkBC_Vec_MC_Port_Ordering_Transactors
        (Tuple2 #(Vector #(16, BC_MC_Client),
                  Vector #(16, BC_MC_Server)));
```

Again, the servers are inward-looking to the app, and the clients implement the outward-looking interfaces. The ordering transactor allows up to 100 requests to be in flight (please see source code for exact size, or if you wish to modify it). To accomplish this, it also needs to use 7 bits of the `rtnctl` tag in memory read-requests, which it shifts in from the LSB side. Thus, if the app actually uses the `rtnctl` field for anything, it should only use 25 bits (=32-7) on the LSB end, to avoid losing information.

D.3.3   Management null transactor

The Management null transactor simply "loops back" the CSR ring. It makes the `csr_31_31_intlv_dis` and `aeid` values available to the app using the following inward-looking interface:

```
────────────────────── BC_Transactors.bsv ──────────────────────
interface BC_Management_App_IFC;
   (* always_ready *)
   method BC_AEId  aeid;

   (* always_ready *)
   method Bool  csr_31_31_intlv_dis;
endinterface
```

D.4   Memory and on-chip networks

Designs for the FPGAs often need various kinds of networks. Examples:

- Memory networks: memory requests from deep within the module hierarchy have to be brought to the top of the hierarchy and routed into the correct MC port. Memory responses have to be routed back from MC ports to the request origination points.

- Data distrubution and collection: tasks sent from the SW side via a memory FIFO must be distributed to several concurrent "worker" modules. Results from these workers must be gathered and sent back via a memory FIFO to the SW sides.

BClib has facilities to help build such networks.

As any textbook on the subject will describe, network implementations typically have trade-offs between latency, bandwidth, and hardware resources. All the networks described in this section are implemented as "ring nets". This is a linear circular buffer as large as the number of inputs or outputs. Each of the input ports connects (in parallel) to one of the buffer slots. Each of the output ports connects (in parallel) to one of the buffer slots. Incoming packets enter in some buffer slot and circulate around the ring until come across their destination port, where they exit. This is a reasonable point in the trade-off space: low hardware resources with some parallism (bandwidth), with not particularly great latency. However, note that these implementations are internal details of the particular modules we provide; you could easily define alternative topology network implementations for exactly the same interfaces.

Please also see [3] for an automated web-based tool to create custom networks on FPGAs (it uses BSV internally).

D.4.1   Memory request-response fat trees: `BC_Mem_Tree.bsv`

The file `BC_Mem_Tree.bsv` provides types and modules to construct "fat tree" memory routing networks. The key interface is:

```
───────────────────── BC_Mem_Tree.bsv ─────────────────────
interface MemTreeNode #(numeric type m, numeric type n);
   method Action init (BC_AEId aeid);
   interface Vector #(n, BC_Mem_Client)  rootside;
   interface Vector #(m, BC_Mem_Server)  leafside;
endinterface
```

which describes a fat tree node with m leaf-side server (request/response) ports and n root-side ports (a conventional tree is just a special case of a fat tree where when $n = 1$).

The package provides a module to construct such a fat tree node for memory requests and responses:

```
───────────────────── BC_Mem_Tree.bsv ─────────────────────
module mkMemTreeNode #(Bool isRoot) (MemTreeNode #(m, n));
```

When `isRoot` is false, it builds an intermediate node for a tree, which uses the uplinks (towards the root) in a random order to balance the traffic load. When `isRoot` is true, it builds a root node for a tree in which packets are routed to a specific uplink node according to the Convey MC interleaving structure, i.e., it directs packets to the correct MC port. The module also automatically handles the response routing. When requests flow up the tree, it automatically records the route as a Convey MC "rtnctl" tag. When a response comes back from an MC with this tag, it is used to route the response down the tree to the corresponding originator.

Using this module, you can construct a custom fat tree that is overlaid on your application module structure.

D.4.2   Generic m-input n-output networks: `RingNet.bsv`

The file `RingNet.bsv` provides generic m-input, n-output network modules.

```
                          ──────── RingNet.bsv ────────
interface RingNet_IFC #(numeric type m, numeric type n, type t_pkt);
    interface Vector #(m, Put #(t_pkt)) inputs;
    interface Vector #(n, Get #(t_pkt)) outputs;


    method Action setDrainState (Bool state);
    method Bool drained;
endinterface
```

`t_pkt` is the type of the packets carried by the network. The first two lines are the m inputs and n outputs. The `setDrainState` method can be used to stop the node from accepting more inputs until all packets within the network have been drained. The `drained` method signals this status.

The following module implements such a network:

```
                          ──────── RingNet.bsv ────────
module mkRingNet
    #(function Bool destination (t_pkt pkt, Bit #(log_n) output_number))
    (RingNet_IFC #(m, n, t_pkt));
```

The `destination` parameter is a function that decides whether a packet `pkt`'s destination is `output_number`. Thus, you can very flexibly choose how packets are routed.

E   Release Notes

E.1   Release 2012.08.A

Original release.

E.2   Release 2013.08.A

New version of `CompletionBuffer2.bsv` in `BClib_src` directory. Replaced the "vector of registers" implementation of `mkCompletionBuffer2` with a BRAM-based implementation, thereby improving place and route, timing closure, etc.

New version of `SpecialFIFOS2.bsv` in `BClib_src` directory. Replaced the "vector of registers" implementation of `mkSizedFIFOF_AlmostFull` with a BRAM-based implementation, thereby improving place and route, timing closure, etc. This is used to implement buffers for memory responses; there are 16 of them for the 8 even-and-odd memory channels.

E.3   Release 2014.12.C

This release is for Convey MX platforms. The older HC version is available as well, for those customers who might be developing for both platforms. The top-level of the distribution is separated into HC and MX directories for this purpose.

The new MX-compatible code uses Convey's new "Type 1" memory interface. The main differences from the older Type 0 memory interface are:

- Store operations now have responses (previously, they were "fire and forget", i.e., there was no response, and the only way to ensure store completion was to do a flush request/response. But that was a heavyweight hammer, encompassing all store ops, whereas now we can ensure completion of individual stores.

- Extensive suite of Atomic memory ops in the hardware. These are fully implemented in the BClib Bluesim simulation environment as well.

- Extensive suite of Full/Empty bit ops in the hardware. While defined in BClib and usable from BClib when running on Convey hardware, these are NOT implemented in the BClib Bluesim simulation environment, since it is not clear how this can be implemented for arbitrary memory locations. (This functionality also seems to be dropped in Convey's next platform, the WX "Wolverine" series).

Dropped the AE-to-AE interfaces completely. Even in previous versions, the functionality was never implemented, but stubs were provided to enable correct interfacing to Convey's PDK. Subsequently, Convey's PDK has made these optional (with `AE_AE_IF=0` in the build Makefile), so we have removed it completely from BClib.

Generated Verilog now directly uses the native Convey PDK-provided positive reset signal. In previous releases, the PDK-provided signal was inverted and fed to the rest of the BSV code. While functionally correct, the presence of an ordinary (not BUFG) inverter in the reset path sometimes caused timing closure problems.

References

[1] Bluespec, Inc. Bluespec^{TM} SystemVerilog Version Reference Guide, 2011.

[2] Bluespec, Inc. Bluespec^{TM} SystemVerilog Version User Guide, 2011.

[3] M. Papamichael. CONNECT: Configurable Network Creation Tool. Free service (for non-commercial use) at http://www.ece.cmu.edu/ mpapamic/connect/ that generates on-chip FPGA networks based on your specification.